

# Lecture 2: Networking Concepts

Oct 3

# Administrivia


- Project #1 out tomorrow
  - Will post on Piazza
  - Due 10/12 (Friday 11:59PM)
- Next M: Guest lecture
- Next W: No lecture, MORE time for project 1


# Recap: The Internet is...

- A federated system
- Of enormous scale
- Dynamic range
- Diversity
- Constantly evolving
- Asynchronous in operation
- Failure prone
- Constrained by what's practical to engineer
- Too complex for theoretical models
- “Working code” needn't mean much
- Performance benchmarks are too narrow

# Today

- What is a network made of?
- How is it shared?
- Sockets
- How do we evaluate a network?







# Cable


- Coaxial copper & fiber
- up to 42.8 Mbps downstream
- up to 30.7 Mbps upstream
- shared broadcast medium

# Hybrid Fiber-Coax










# Ethernet

- Twisted pair copper
- 100 Mbps, 1 Gbps, 10 Gbps (each direction)


# & more

- Cellular (smart phones)
- Satellite (remote areas)
- Fiber to the Home (home)
- Optical carrier (Internet backbone)




# Today

- What is a network made of?
- How is it shared?
- Sockets
- How do we evaluate a network?







How do we scale a network to many end-systems?



$O(n^2)$  links




per-node capacity scales as  $1/n$







How do we scale a network to many end-systems?




Switched networks enable efficient scaling!

# Two approaches to sharing



- Reservations
- On demand

# Example: Three sources w/ “bursty” traffic

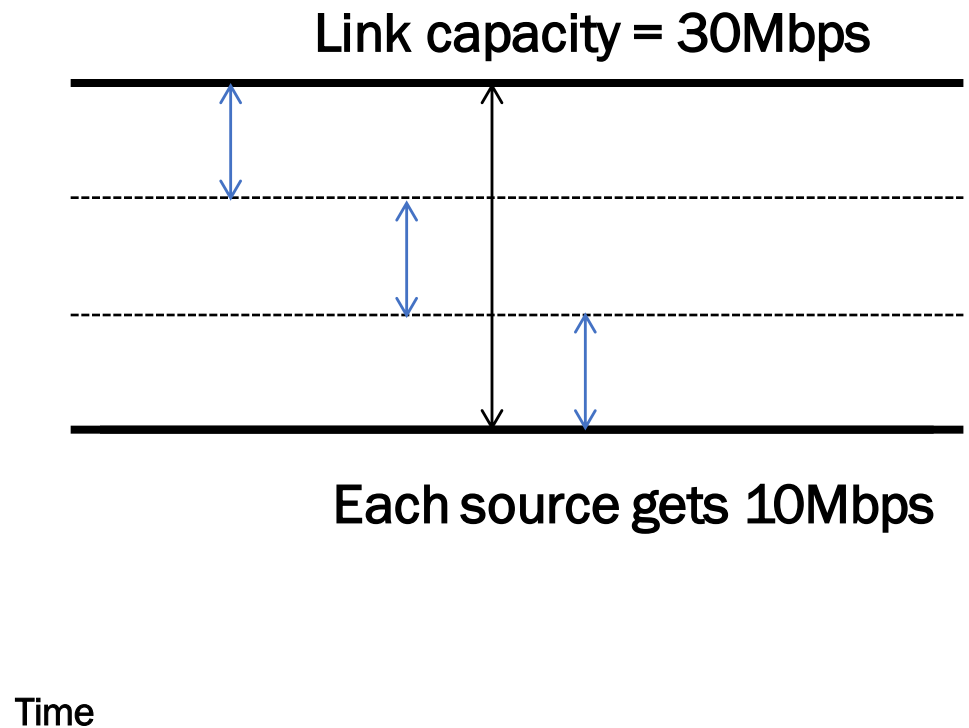


Time



# Intuition: reservations



**Frequent overloading**




# Intuition: on demand



Time



**No overloading**



Switching on-demand exploits *statistical multiplexing* better than reservations

- Sharing using the statistics of demand
- Good for bursty traffic (average  $\ll$  peak demand)
- Similar to insurance, with the same failure mode

# Intuition: on demand



**What do we do under overload?**

# Statistical multiplexing is a recurrent theme in computer science

- Phone network rather than dedicated lines
  - ancient history
- Packet switching rather than circuits
  - today's lecture
- Cloud computing
  - shared vs. dedicated machines



# Two approaches to sharing

- Reservations
- On demand

**How are these implemented?**

# Two approaches to sharing


- Reservations → circuit switching
- On demand → packet switching

How are these implemented?

# Two approaches to sharing


- Packet switching
  - packets treated on demand
  - admission control: per packet
- Circuit switching
  - resources reserved per active "connection"
  - admission control: per connection
- A hybrid: virtual circuits
  - emulating circuit switching with packets (see text)

# Circuit Switching



- (1) **src** sends a reservation request to **dst**
- (2) Switches “establish a circuit”
- (3) **src** starts sending data
- (4) **src** sends a “teardown circuit” message



# Circuit Switching




Reservation establishes a “circuit” within a switch

# Many kinds of “circuits”


- Time division multiplexing
  - divide time in **time slots**
  - separate time slot per circuit
- Frequency division multiplexing
  - divide frequency spectrum in **frequency bands**
  - separate frequency band per circuit



# Timing in Circuit Switching




# Timing in Circuit Switching







# Timing in Circuit Switching





# Timing in Circuit Switching





# Timing in Circuit Switching





# Timing in Circuit Switching




# Timing in Circuit Switching




# Timing in Circuit Switching



# Timing in Circuit Switching




# Timing in Circuit Switching





# Circuit Switching




Circuit switching doesn't "route around trouble"

# Circuit Switching


- Pros
  - predictable performance
  - simple/fast switching (once circuit established)
- Cons
  - complexity of circuit setup/teardown
  - inefficient when traffic is bursty
  - circuit setup adds delay
  - switch fails → its circuit(s) fails

# Packet switching



Each packet contains destination (**dst**)  
Each packet treated independently

# Packet switching



Each packet contains destination (**dst**)

Each packet treated independently

With buffers to absorb transient overloads

# Packet Switching

- Pros
  - efficient use of network resources
  - simpler to implement
  - robust: can “route around trouble”
- Cons
  - unpredictable performance
  - requires buffer management and congestion control

- *On-demand or reserve?*

# Today

- What is a network made of?
- How is it shared?
  - **will assume packet switching from now on**
- Sockets
- How do we evaluate a network?

# Time Travel to Sockets

The background is a vibrant, abstract composition of light trails and a central glowing sphere. The light trails are composed of numerous thin, parallel lines in shades of purple, blue, orange, and yellow, creating a sense of rapid motion and depth. These trails converge towards a central point where a bright, multi-colored sphere is located. The sphere has a core of intense yellow and orange, surrounded by layers of purple, blue, and pink. The overall effect is reminiscent of a wormhole or a high-speed journey through space and time.



# Quick tutorial of Sockets

<http://www.cs.rpi.edu/~moorthy/Courses/os98/Pgms/socket.html>

# Sockets

- Sockets
  - Basic API for programming network applications
  - Should have been covered (some) in CS 15400



# Overview: Communication via Sockets

- Client

1. Create a socket with the **socket()** system call
2. Connect the socket to the address of the server using the **connect()** system call
3. Send and receive data. The simplest is to use the **read()** and **write()** system calls.

- Server

1. Create a socket with the **socket()** system call
2. Bind the socket to an address using the **bind()** system call.
3. Listen for connections with the **listen()** system call
4. Accept a connection with the **accept()** system call (blocks until a client connects with the server)
5. Send and receive data

# Application Programming Interface

- Application interfaces to protocols can be
  - *Loosely specified*
  - *Specified precisely*
- An API is a set of function signatures
- It is an interface that may have multiple implementations
  - *System calls vs. libraries*
  - *Unix, Windows*
- An API can define dedicated abstractions/functions or reuse existing ones

# Socket API

- API used to access the network
- A socket is an abstraction of a communication end-point
  - *Can be used to access different network protocols (not only TCP/IP)*
- A socket's characteristics are determined by calling specific functions
- A socket can be accessed as a file
  - *Some of the I/O API calls are similar (e.g., read, write, close)*

# Creating a Socket

- `int socket(int domain, int type, int protocol)`
  - **domain**: `PF_UNIX, PF_INET, PF_INET6`
  - **type**: `SOCK_STREAM, SOCK_DGRAM`
  - protocol: normally 0, can be other values if there are multiple protocols available (`IPPROTO_TCP, IPPROTO_UDP`)
  - Return value: the file descriptor of the socket or -1 in case of errors
- UDP and TCP accessed using the same primitives but semantics are different

Protocol 0 = IP  
UDP or TCP chosen according to type

# Socket Domain & Type

- Two address domains:
  - **Unix** domain: when two processes sharing a common file system to communicate with each other
  - **Internet** domain: when two processes running on any two hosts on the Internet communicate.
- Two socket types:
  - **Stream**: continuously stream of data (TCP)
  - **Datagram**: message as a block (UDP)

# Socket Addresses

Really just a wrapper  
that provides format  
of data

```
struct sockaddr {  
    u_short sa_family; /* family */  
    char sa_data[14]; /* address data */  
};
```

```
struct sockaddr_in { /* a TCP endpoint */  
    u_int16_t sin_family; /* address family: AF_INET */  
    u_int16_t sin_port; /* port in network byte order */  
    struct in_addr sin_addr; /* internet address */  
};
```

Internet namespace,  
includes address and port

```
struct in_addr {  
    u_int32_t s_addr; /* address in network byte order */  
};
```

Headers defined  
In sys/socket.h



# Socket Descriptors

- Similar to file descriptors
  - *Internal data structure specifies*
    - *Protocol used (or family – PF\_INET)*
    - *Type of service (connectionless or connection-oriented)*
    - *IP addresses involved*
    - *Ports involved*

**Essentially a “handle”  
on the socket, which can  
be used to do things like  
“accept.”**

- Per-process resource

# Overview: Communication via Sockets

- Client

1. Create a socket with the **socket()** system call
2. Connect the socket to the address of the server using the **connect()** system call
3. Send and receive data. The simplest is to use the **read()** and **write()** system calls.

- Server

1. Create a socket with the **socket()** system call
2. Bind the socket to an address using the **bind()** system call.
3. Listen for connections with the **listen()** system call
4. Accept a connection with the **accept()** system call (blocks until a client connects with the server)
5. Send and receive data

# Creating Connections

- Clients use **connect()** to open a connection to a specific IP address/port
- **int connect(int sockfd, const struct sockaddr \*serv\_addr, socklen\_t addrlen)**
  - **sockfd**: socket descriptor
  - **serv\_addr**: destination address
  - **addrlen**: size of the structure
  - Return value: 0 in case of success, -1 in case of errors

# Connecting with TCP and UDP

- TCP

- ***connect ()** starts the three-way handshake*

- UDP

- *Nothing really happens, but the socket can only be used to send/receive datagrams to/from the specified address*

# Overview: Communication via Sockets

- Client

1. Create a socket with the **socket()** system call
2. Connect the socket to the address of the server using the **connect()** system call
3. Send and receive data. The simplest is to use the **read()** and **write()** system calls.

- Server

1. Create a socket with the **socket()** system call
2. Bind the socket to an address using the **bind()** system call.
3. Listen for connections with the **listen()** system call
4. Accept a connection with the **accept()** system call (blocks until a client connects with the server)
5. Send and receive data

# Binding Sockets to Addresses

- A socket can be bound to a specific IP address and port
- **`int bind(int sockfd, struct sockaddr *my_addr, socklen_t addrlen)`**
  - ***sockfd**: socket descriptor*
  - ***my\_addr**: socket address (usually of `sockaddr_in` type)*
  - ***addrlen**: address structure length*
  - *Return value: 0 in case of success, -1 in case of error*
- Same semantics for TCP and UDP

# Passive Connection-Oriented Sockets

- If a socket is used to receive TCP connections, it is necessary to bind it to an address and to specify that the socket is passive
- **int listen(int s, int backlog)**
  - **s**: socket descriptor
  - **backlog**: maximum length of the queue of pending connections (used to be the number of open/half-open connections, now only the open ones that are ready to be accepted are counted)
  - Return value: 0 in case of success, -1 in case of error

Why don't we bind for active sockets?

# Accepting Connections

- **int accept(int s, struct sockaddr \*addr, socklen\_t \*addrlen)**
  - **s**: socket descriptor
  - **addr**: structure that will be filled with the parameters of the **client** (maybe NULL)
  - **addrlen**: length of the structure in input, it is filled with the actual size of the data returned
  - Return value: a new socket file descriptor in case of success, -1 in case of errors



# Accepting TCP Connections

- A call to **accept()** blocks the caller until a request is sent
- When a connection is made, **accept()** returns a new socket, associated with a specific client (virtual circuit)
- A virtual circuit is identified by:  
**<srcIP, srcPort, dstIP, dstPort>**
- There are no two identical virtual circuits at one time in the whole Internet

# Sending Data On A Connected Socket

- `int send(int s, const void *msg, size_t len, int flags)`
  - `s`: connected socket
  - `msg`: message
  - `len`: size of message
  - `flags`: e.g., `MSG_DONTWAIT`
  - Return value: # of characters sent, or -1 in case of errors

# Receiving Data On A Connected Socket

- `int recv(int s, void *buf, size_t len, int flags)`
  - `s`: connected socket
  - `buf`: buffer to store the message
  - `len`: max size of message in input, actual size of message in output
  - `flags`: e.g., `MSG_PEEK` allows one to peek at a message without removing it from the incoming data queue
  - Return value: the number of characters read, or -1 in case of errors
- This call will block the process if there is no data

# Closing A Socket

- **int close(int fd) ;**
- **fd**: socket descriptor
- Return value: 0 in case of success, -1 in case of errors
  
- **int shutdown(int s, int how)**
- Can be used to close a socket partially
- **s**: connected socket
- **how**:
  - **SHUT\_RD**, further receptions are disallowed
  - **SHUT\_WR**, further transmissions are disallowed
  - **SHUT\_RDWR**, further receptions and transmissions are disallowed
- Return value: 0 in case of success, -1 in case of errors

# Closing Semantics

- TCP
  - Close: FIN/ACK in both directions
  - Shutdown: FIN/ACK in one direction
- UDP
  - Close: Don't send anything. Just deallocate structure

# Sending/Receiving Data on a Disconnected Socket (for UDP)

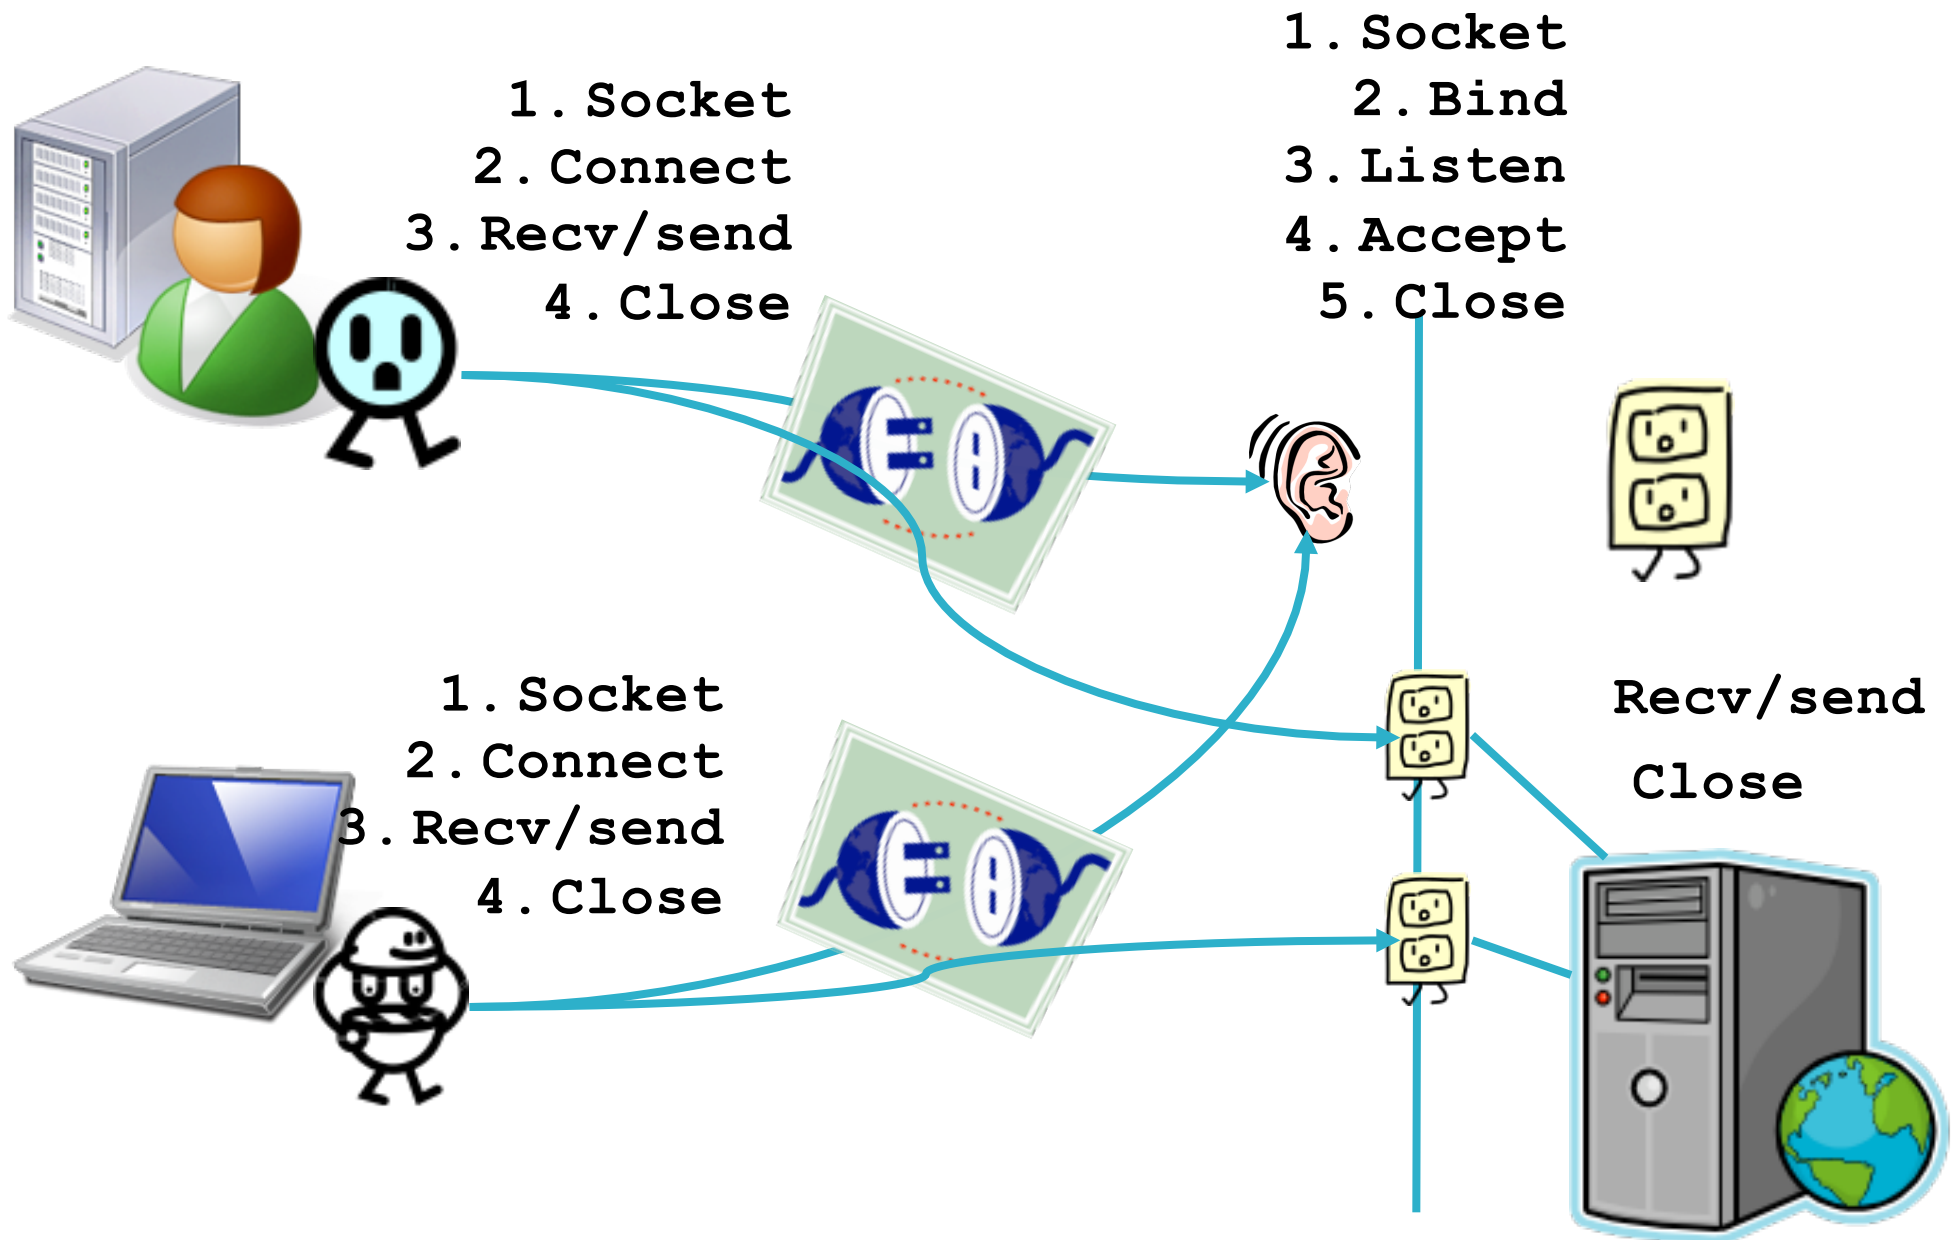
UDP does not require `accept()`. So it need source or destination address information in order to receive or send data.

- `int sendto(int s, const void *msg, size_t len, int flags, const struct sockaddr *to, socklen_t tolen)`
- `int recvfrom(int s, void *buf, size_t len, int flags, struct sockaddr *from, socklen_t *fromlen)`

# Data Exchange in TCP and UDP

- TCP
  - recv/read will return a chunk of data
  - Not necessarily the data sent by means of a single send/write operation on the other side
  - THIS IS IMPORTANT! (semantics of a byte stream, not datagrams)
- UDP
  - recv/read will always return a datagram
  - If message size > buffer, fills buffer + discards rest

# TCP Client and Server





```
/* A simple server in the internet domain using TCP
   The port number is passed as an argument */
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

void error(char *msg)
{
    perror(msg);
    exit(1);
}

int main(int argc, char *argv[])
{
    int sockfd, newsockfd, portno, clilen;
    char buffer[256];
    struct sockaddr_in serv_addr, cli_addr;
    int n;
    if (argc < 2) {
        fprintf(stderr,"ERROR, no port provided\n");
        exit(1);
    }
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");
    bzero((char *) &serv_addr, sizeof(serv_addr));
    portno = atoi(argv[1]);
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = INADDR_ANY;
    serv_addr.sin_port = htons(portno);
    if (bind(sockfd, (struct sockaddr *) &serv_addr,
        sizeof(serv_addr)) < 0)
        error("ERROR on binding");
    listen(sockfd,5);
    clilen = sizeof(cli_addr);
    newsockfd = accept(sockfd, (struct sockaddr *) &cli_addr, &clilen);
    if (newsockfd < 0)
        error("ERROR on accept");
    bzero(buffer,256);
    n = read(newsockfd,buffer,255);
    if (n < 0) error("ERROR reading from socket");
    printf("Here is the message: %s\n",buffer);
    n = write(newsockfd,"I got your message",18);
    if (n < 0) error("ERROR writing to socket");
    return 0;
}
```

```

#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <netdb.h>

```

<http://www.cs.rpi.edu/~moorthy/Courses/os98/Pgm/s/socket.html>

```

void error(char *msg)
{
    perror(msg);
    exit(0);
}

int main(int argc, char *argv[])
{
    int sockfd, portno, n;

    struct sockaddr_in serv_addr;
    struct hostent *server;

    char buffer[256];
    if (argc < 3) {
        fprintf(stderr,"usage %s hostname port\n", argv[0]);
        exit(0);
    }
    portno = atoi(argv[2]);
    sockfd = socket(AF_INET, SOCK_STREAM, 0);
    if (sockfd < 0)
        error("ERROR opening socket");
    server = gethostbyname(argv[1]);
    if (server == NULL) {
        fprintf(stderr,"ERROR, no such host\n");
        exit(0);
    }
    bzero((char *) &serv_addr, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    bcopy((char *)server->h_addr,
        (char *)&serv_addr.sin_addr.s_addr,
        server->h_length);
    serv_addr.sin_port = htons(portno);
    if (connect(sockfd,(struct sockaddr *)&serv_addr,sizeof(serv_addr)) < 0)
        error("ERROR connecting");
    printf("Please enter the message: ");
    bzero(buffer,256);
    fgets(buffer,255,stdin);
    n = write(sockfd,buffer,strlen(buffer));
    if (n < 0)
        error("ERROR writing to socket");
    bzero(buffer,256);
    n = read(sockfd,buffer,255);
    if (n < 0)
        error("ERROR reading from socket");
    printf("%s\n",buffer);
    return 0;
}

```