

IP Routing Algorithms

Lecture 4

CEO



Aide



FedEx



Fedex Envelope (FE)

CEO




Aide



FedEx

What is a Protocol?

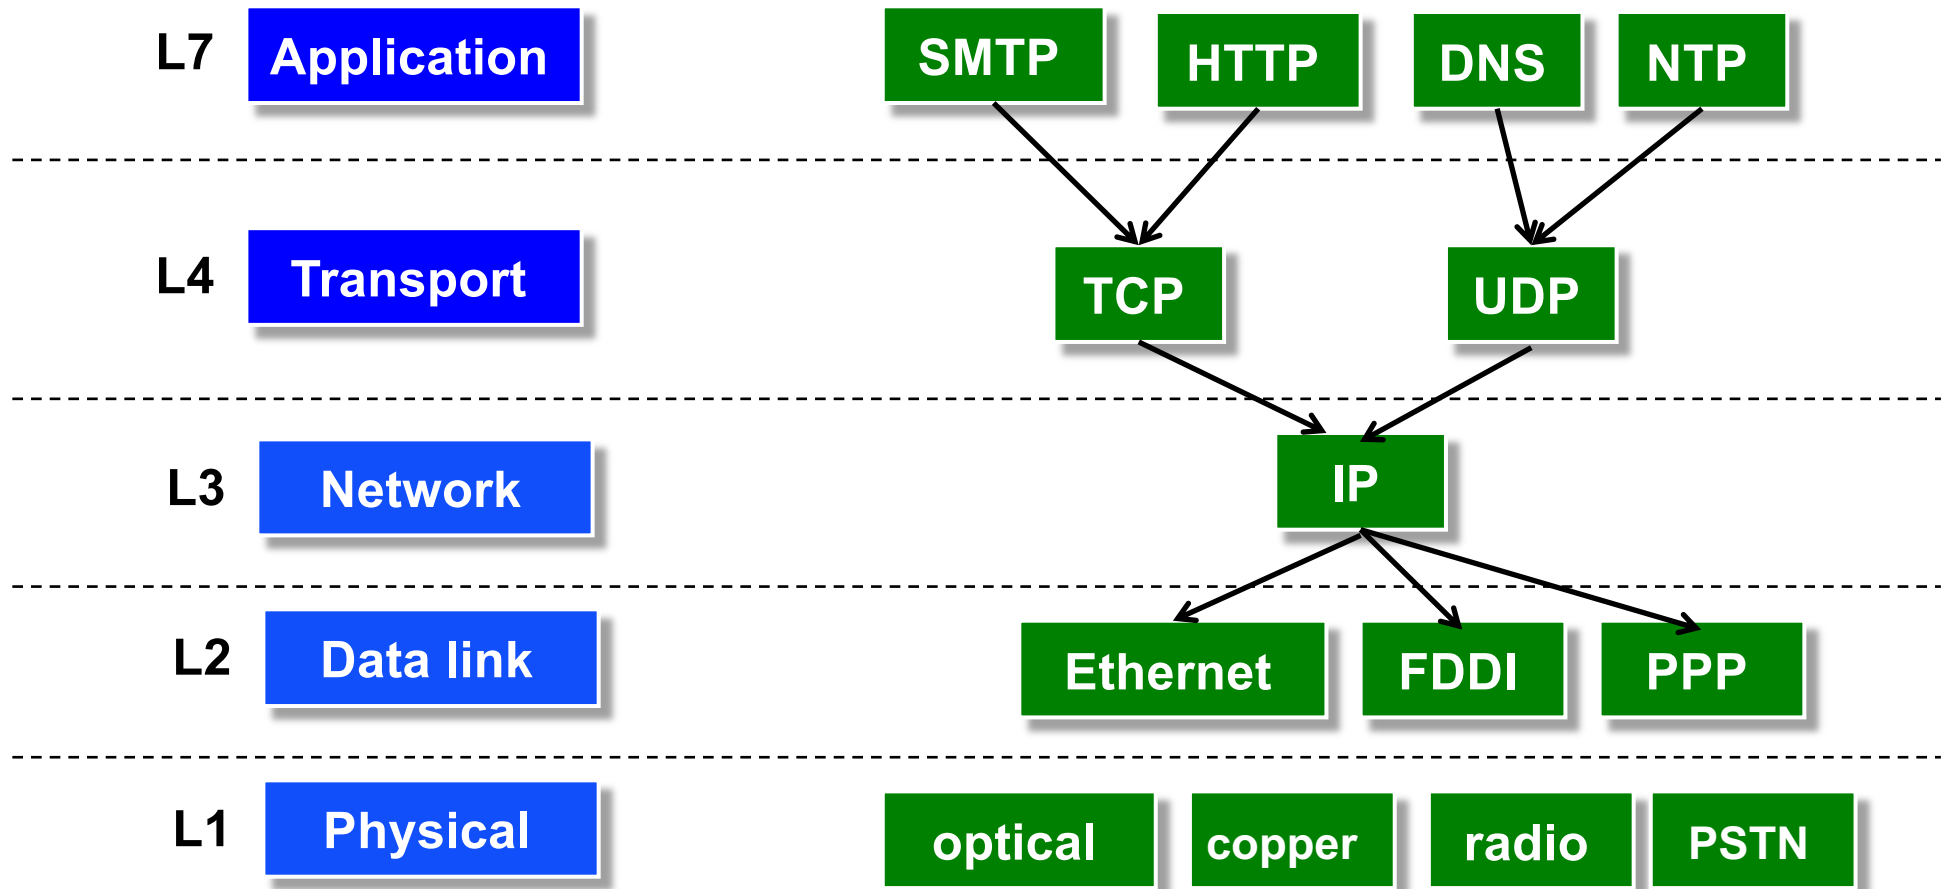
- An agreement between parties on how to communicate
- Defines the syntax of communication
 - header → instructions for how to process the payload
 - Each protocol defines the format of its packet **headers**
 - *e.g. “the first 32 bits carry the destination address”*



What is a Protocol?



- An agreement between parties on how to communicate
- Defines the syntax of communication
- And semantics
 - “first a hello, then a request...”
 - we’ll study many protocols later in the quarter
- Protocols exist at many levels, hardware and software
 - defined by a variety of standards bodies (IETF, IEEE, ITU)

Protocols at different layers



NOTE: just one network-layer protocol!

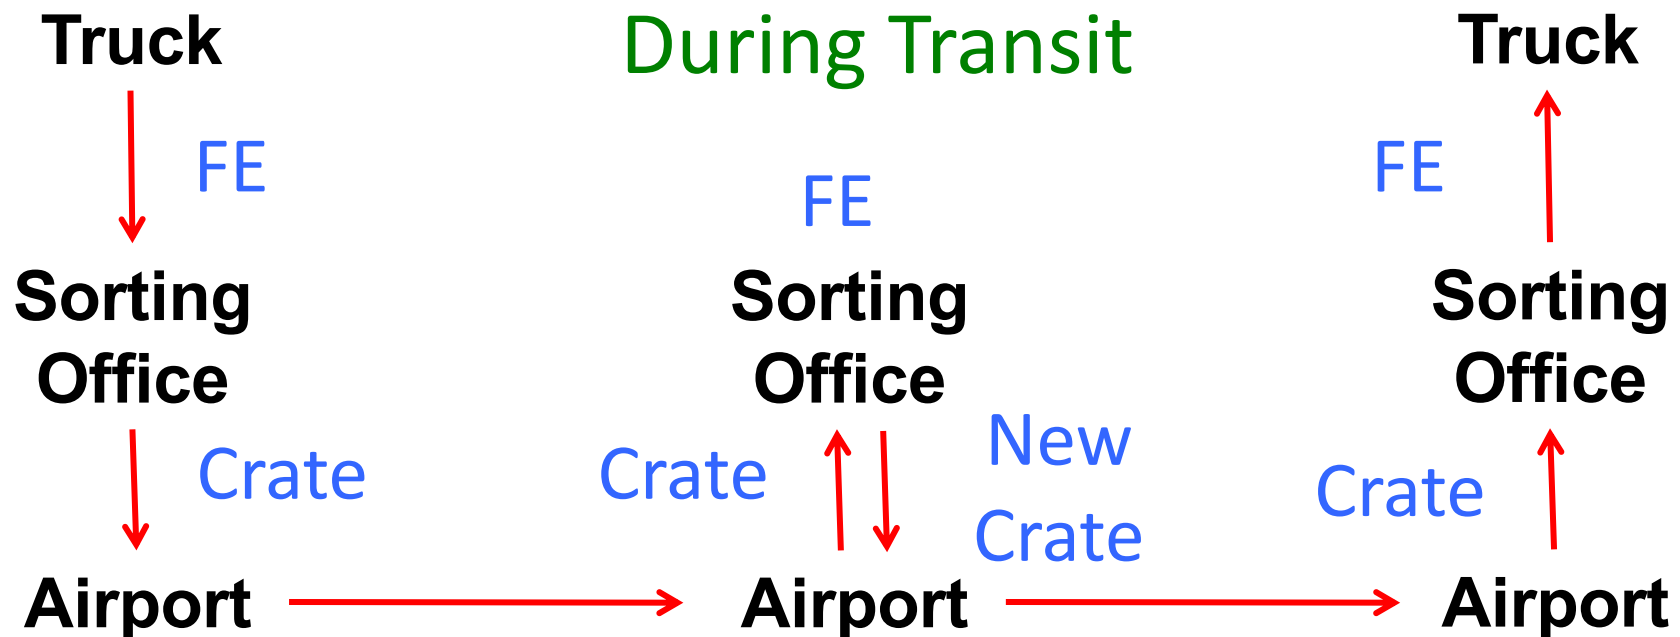
Layer Encapsulation: Protocol Headers




The Path Through FedEx

Higher “Stack”
at Ends

Partial “Stack”
During Transit



Deepest Packaging (Envelope+FE+Crate)
at the Lowest Level of Transport



What gets implemented where?

What gets implemented at the end systems?

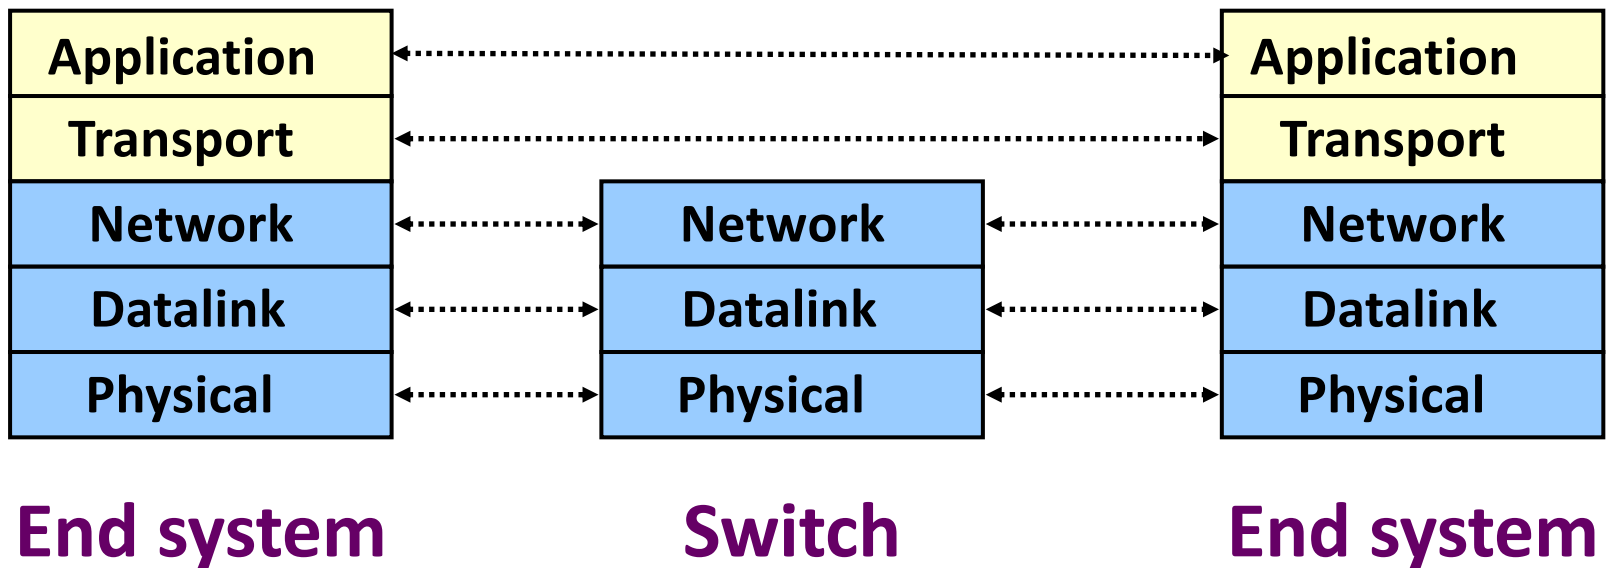
- Bits arrive on wire, must make it up to application
- Therefore, all layers must exist at host!

What gets implemented in the network?


- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across links and local networks → datalink layer (L2)
- Packets must be delivered between networks for global delivery → network layer (L3)
- The network does not support reliable delivery
 - Transport layer (and above) not supported

Simple Diagram


- Lower three layers implemented everywhere
- Top two layers implemented only at hosts



A closer look: network



A closer look: network



What gets implemented in the network?

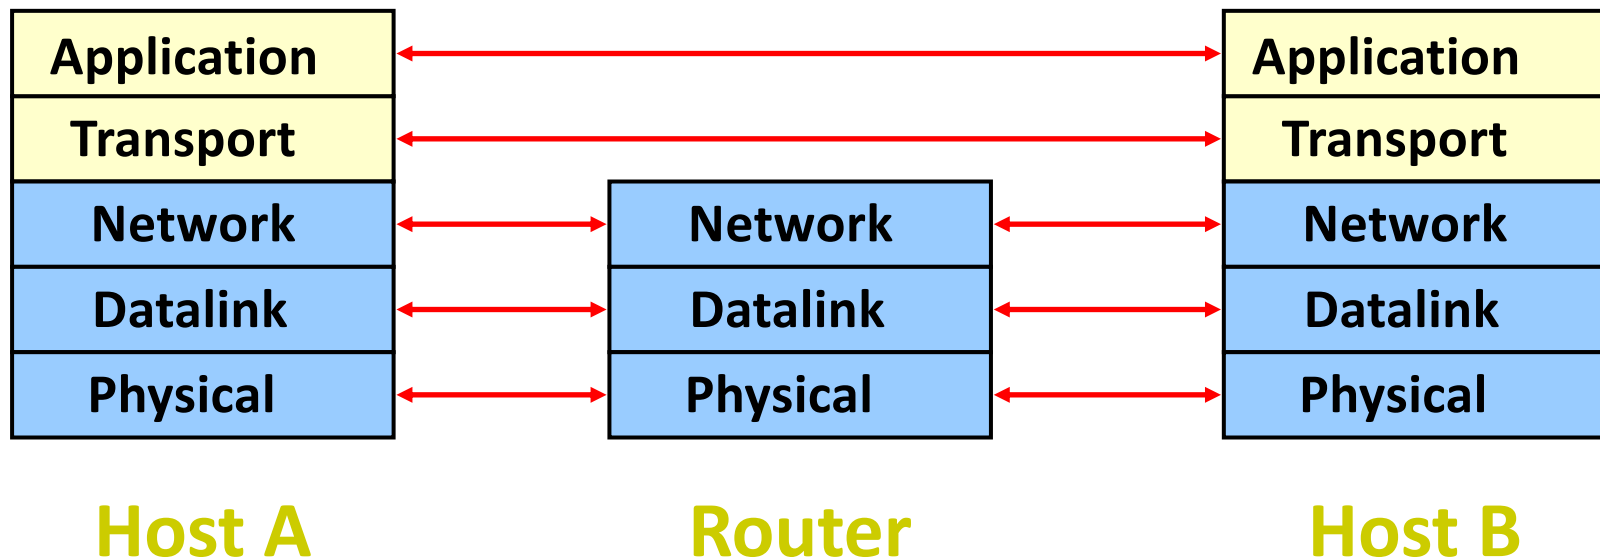
- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across links and local networks → datalink layer (L2)
- Packets must be delivered between networks for global delivery → network layer (L3)
- Hence:
 - switches: implement physical and datalink layers (L1, L2)
 - routers: implement physical, datalink, network layers (L1, L2, L3)

Switches vs. Routers

- Switches do what routers do but don't participate in global delivery, just local delivery
 - switches only need to support L1, L2
 - routers support L1-L3
- Won't focus on the router/switch distinction
 - when I say switch, I almost always mean router
 - almost all boxes support network layer these days

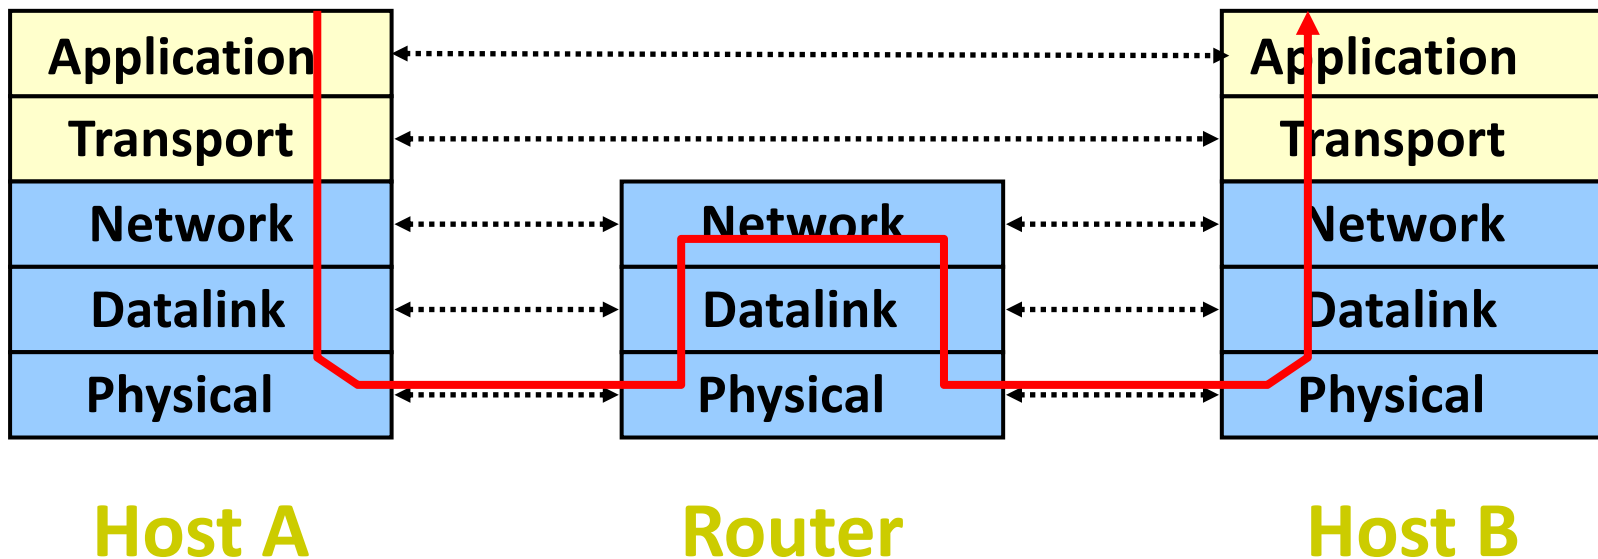
Logical Communication

- Layers interacts with peer's corresponding layer

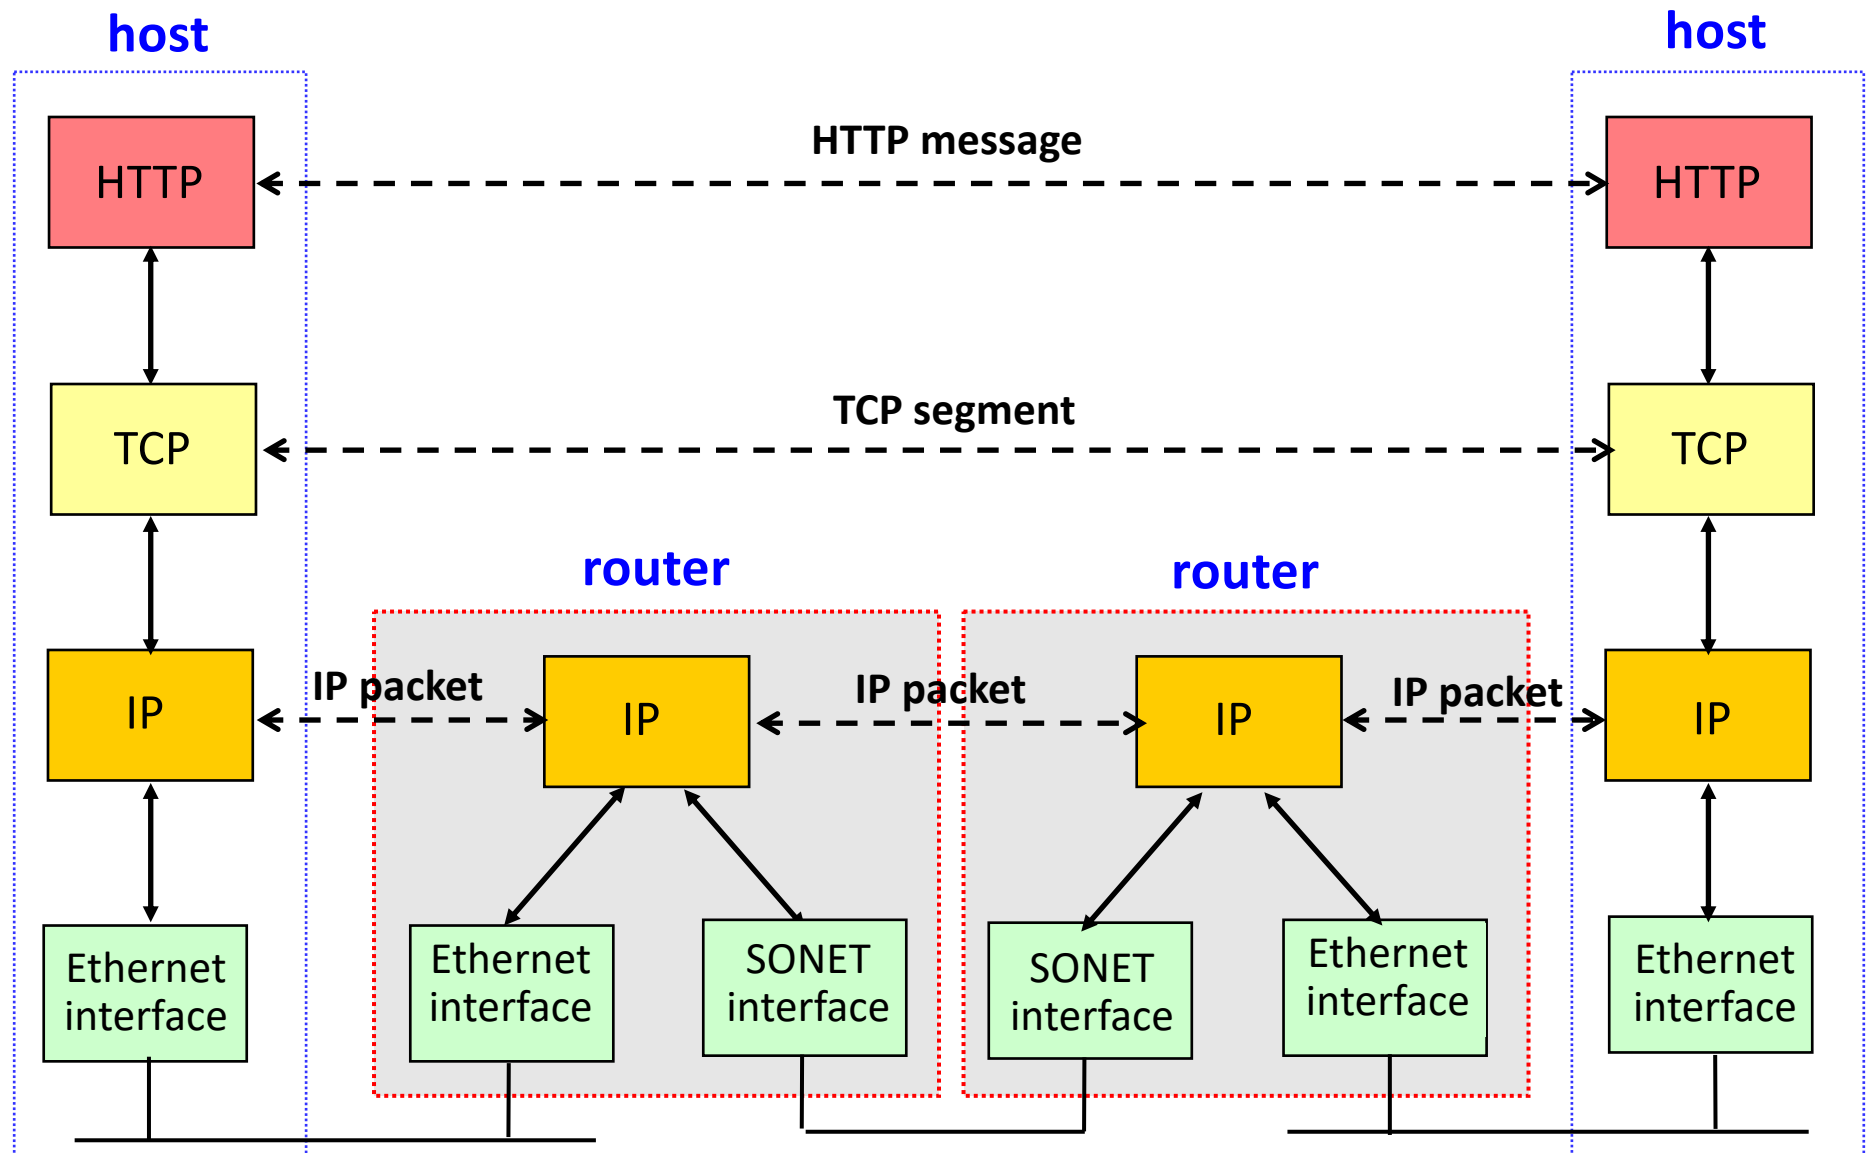


Physical Communication

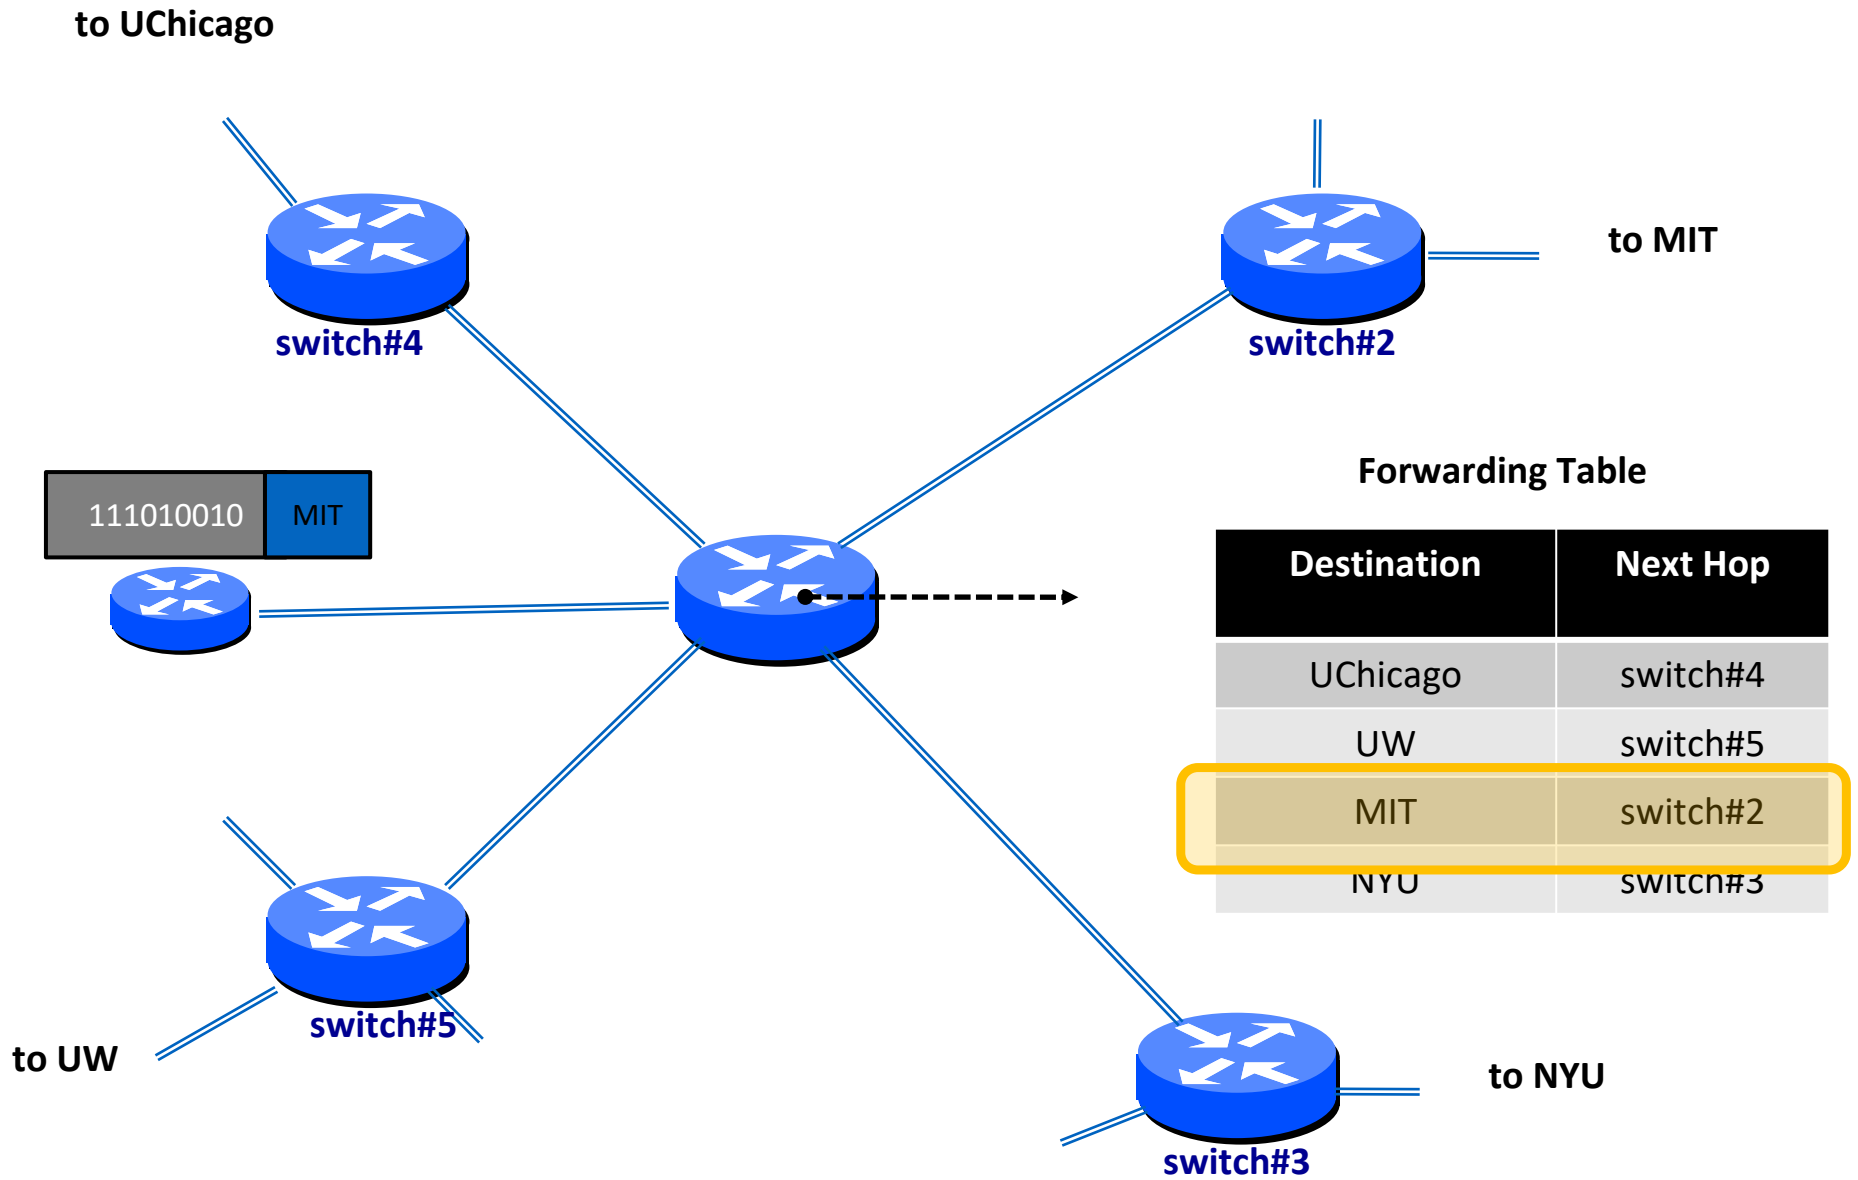
- Communication goes down to physical network
- Then up to relevant layer



A Protocol-Centric Diagram



The Network Layer



Three topics

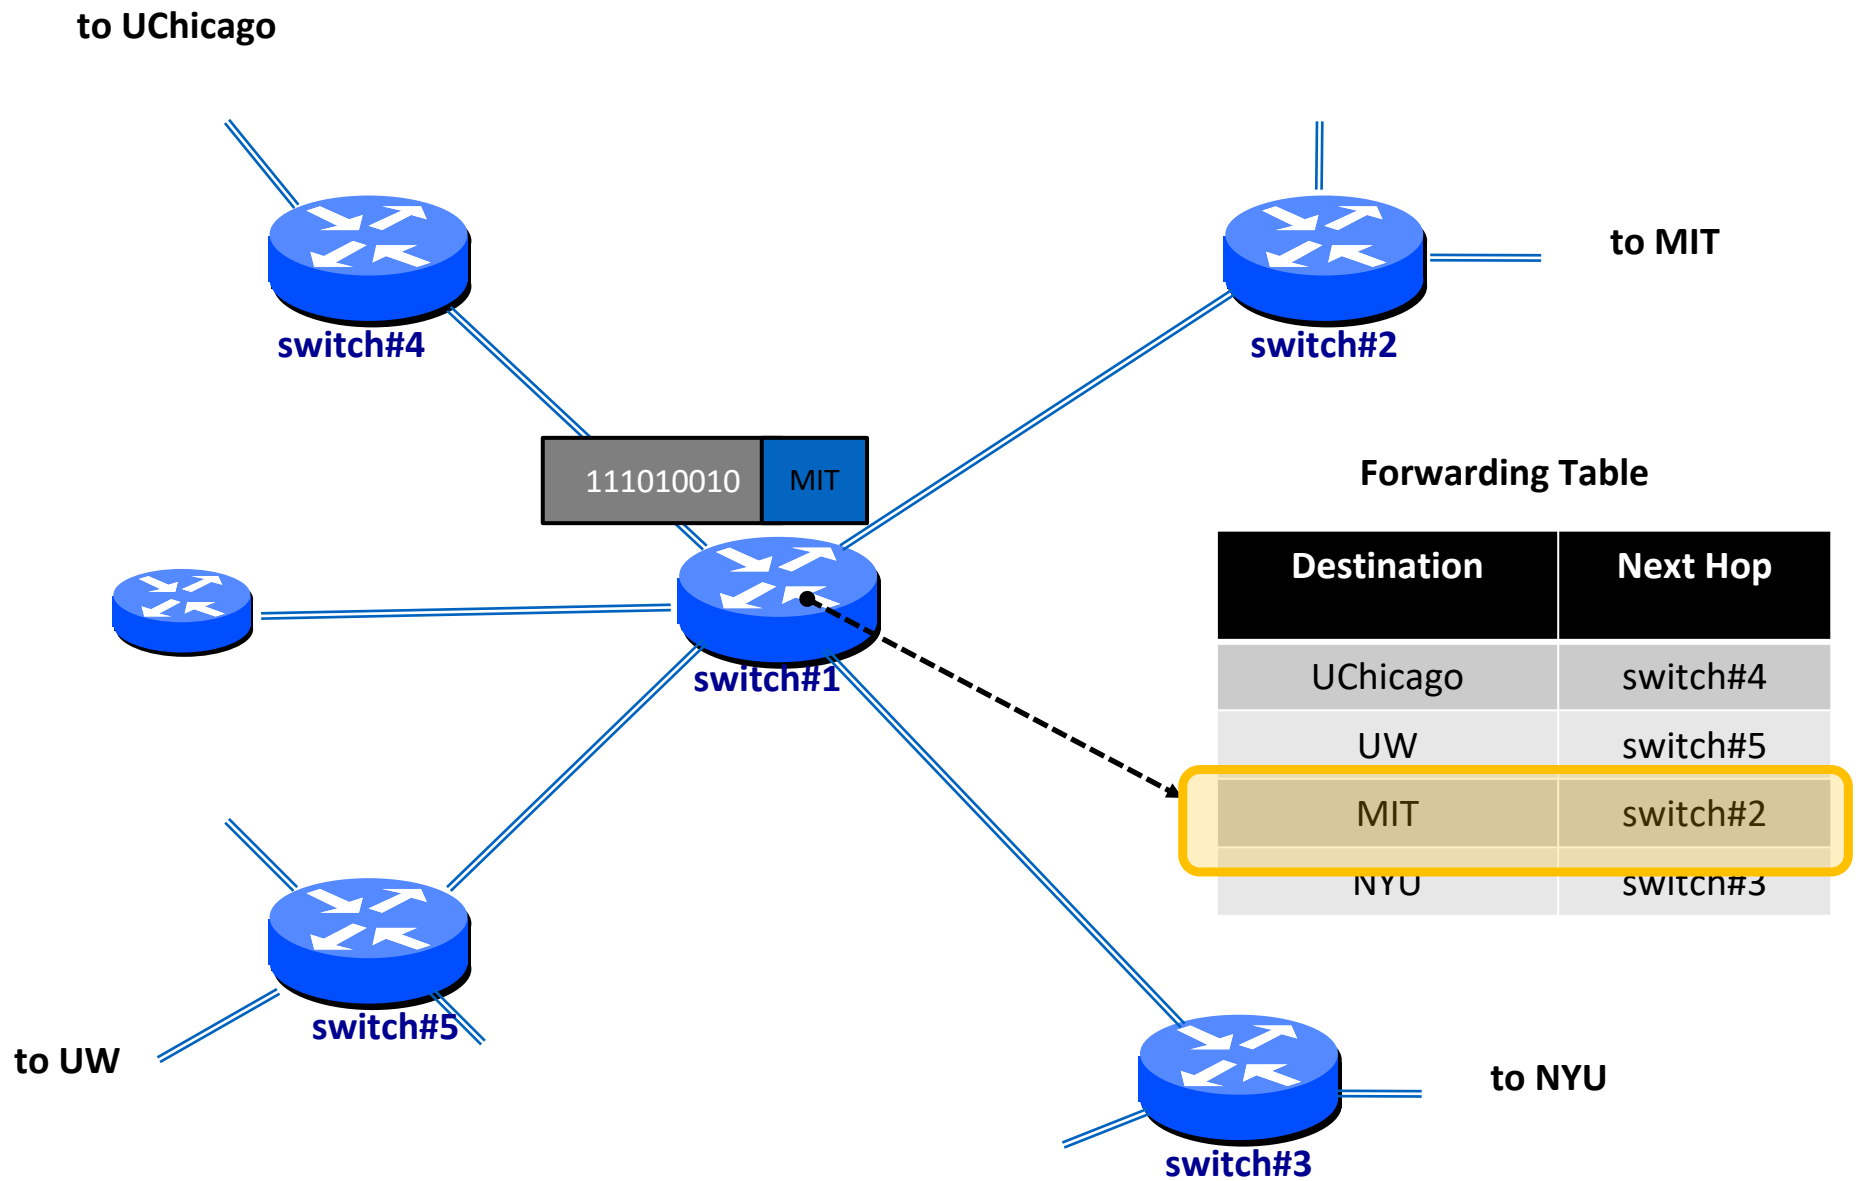
- Addressing
- Forwarding
- Routing

Addressing (for now)

- Assume each end-system has a unique address
- No particular structure to these addresses
- Will cover IP addresses later in the course

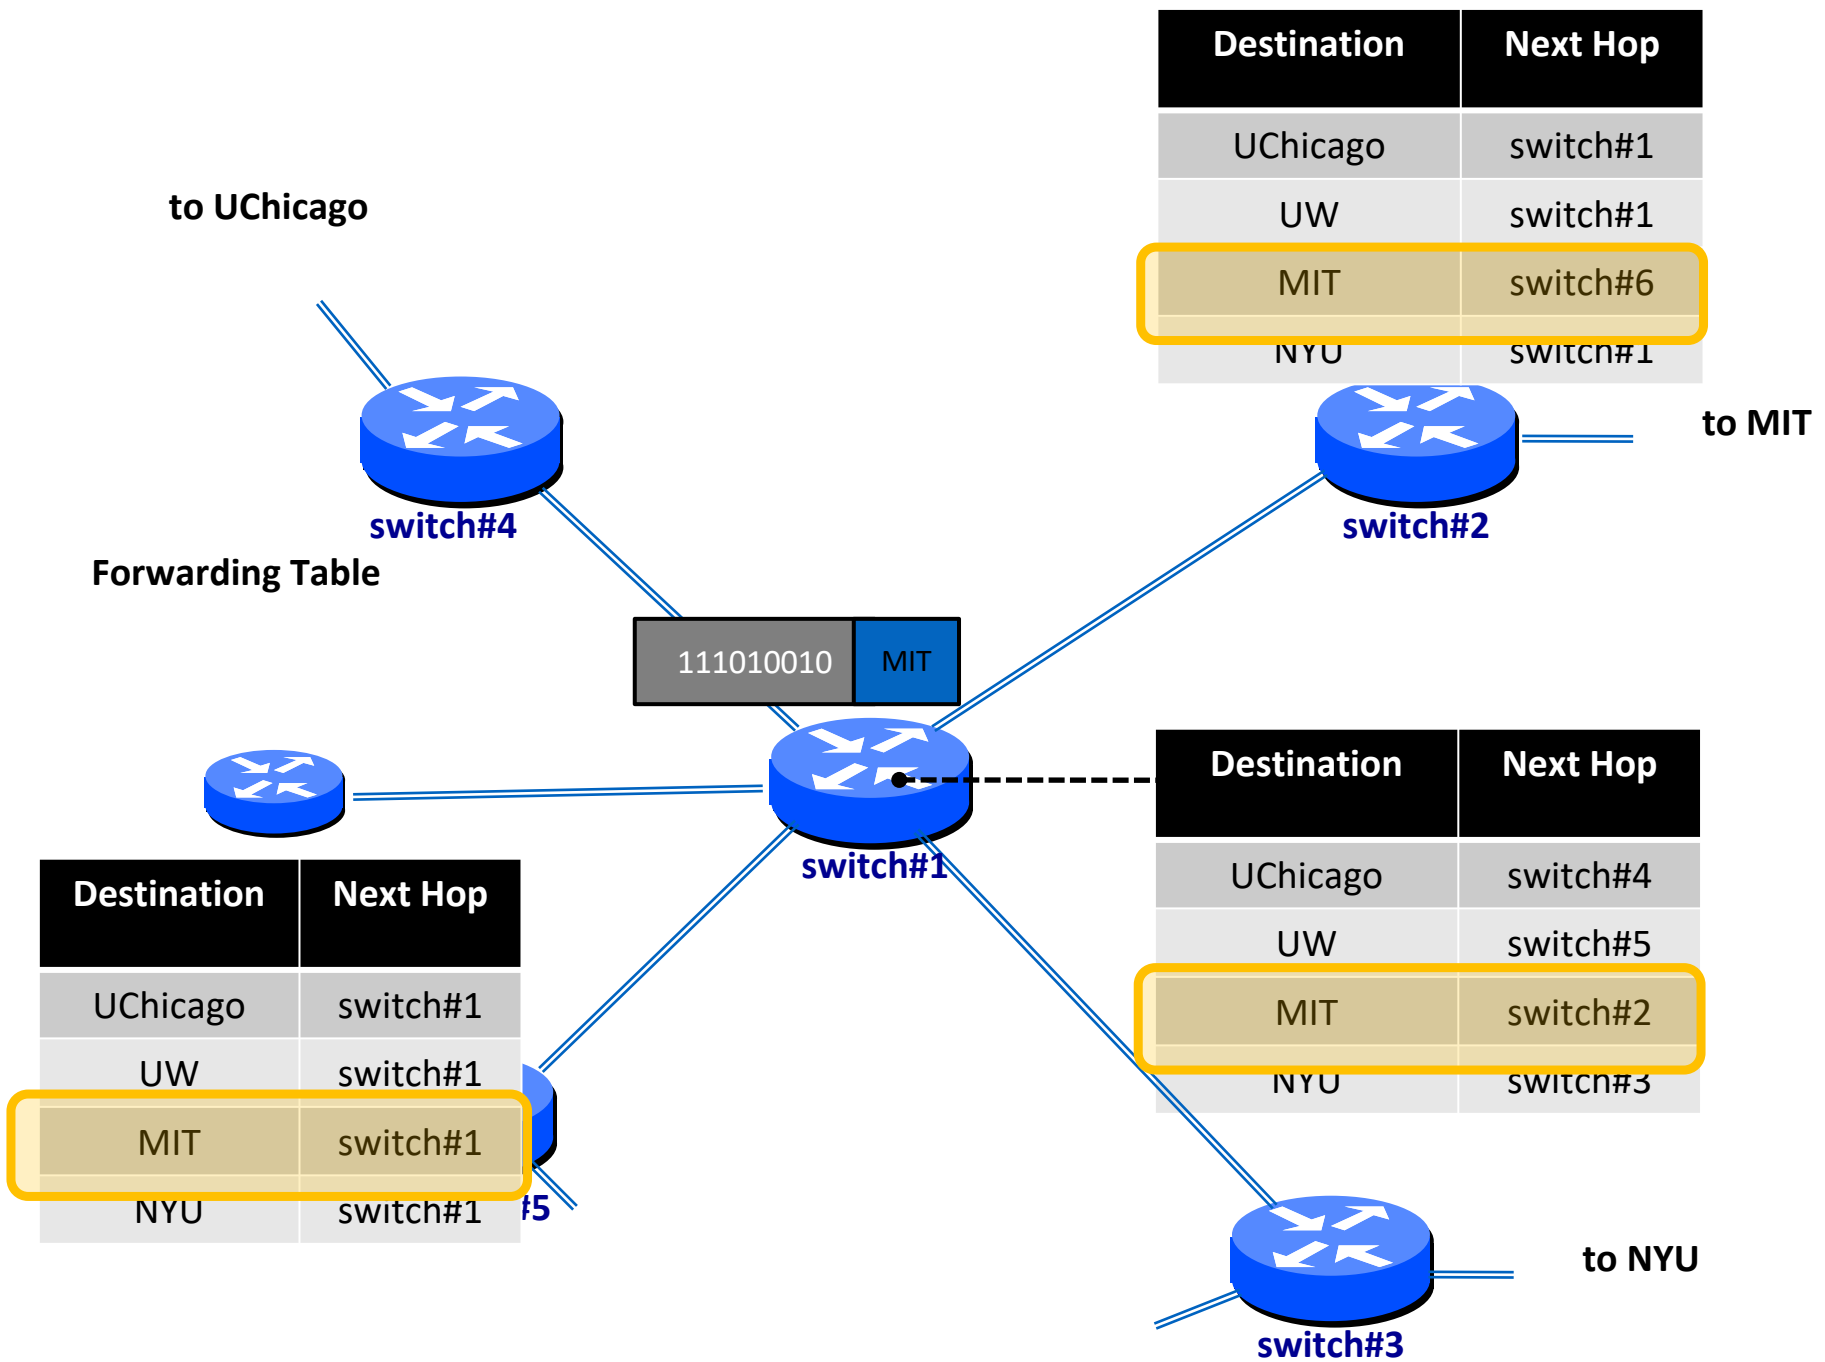
Forwarding

- ▶ **Local** router process that determines the output link (a.k.a “next hop”) for each packet
- ▶ How
 - *read address from packet’s network layer header*
 - *search forwarding table*



Routing

- ▶ **Network-wide** process that determines the content of forwarding tables
 - *determines the end-to-end path for each destination*



Routing

- ▶ **Network-wide** process that determines the content of forwarding tables
 - *determines the end-to-end path for each destination*
- ▶ How
 - *coming up soon*

Forwarding vs. Routing

- ▶ Forwarding: “data plane”
 - *Directing one data packet*
 - *Each router using local routing state*
- ▶ Routing: “control plane”
 - *Computing the forwarding tables that guide packets*
 - *Jointly computed by routers using a distributed algorithm*
- ▶ Very different timescales!

Routing Fundamentals

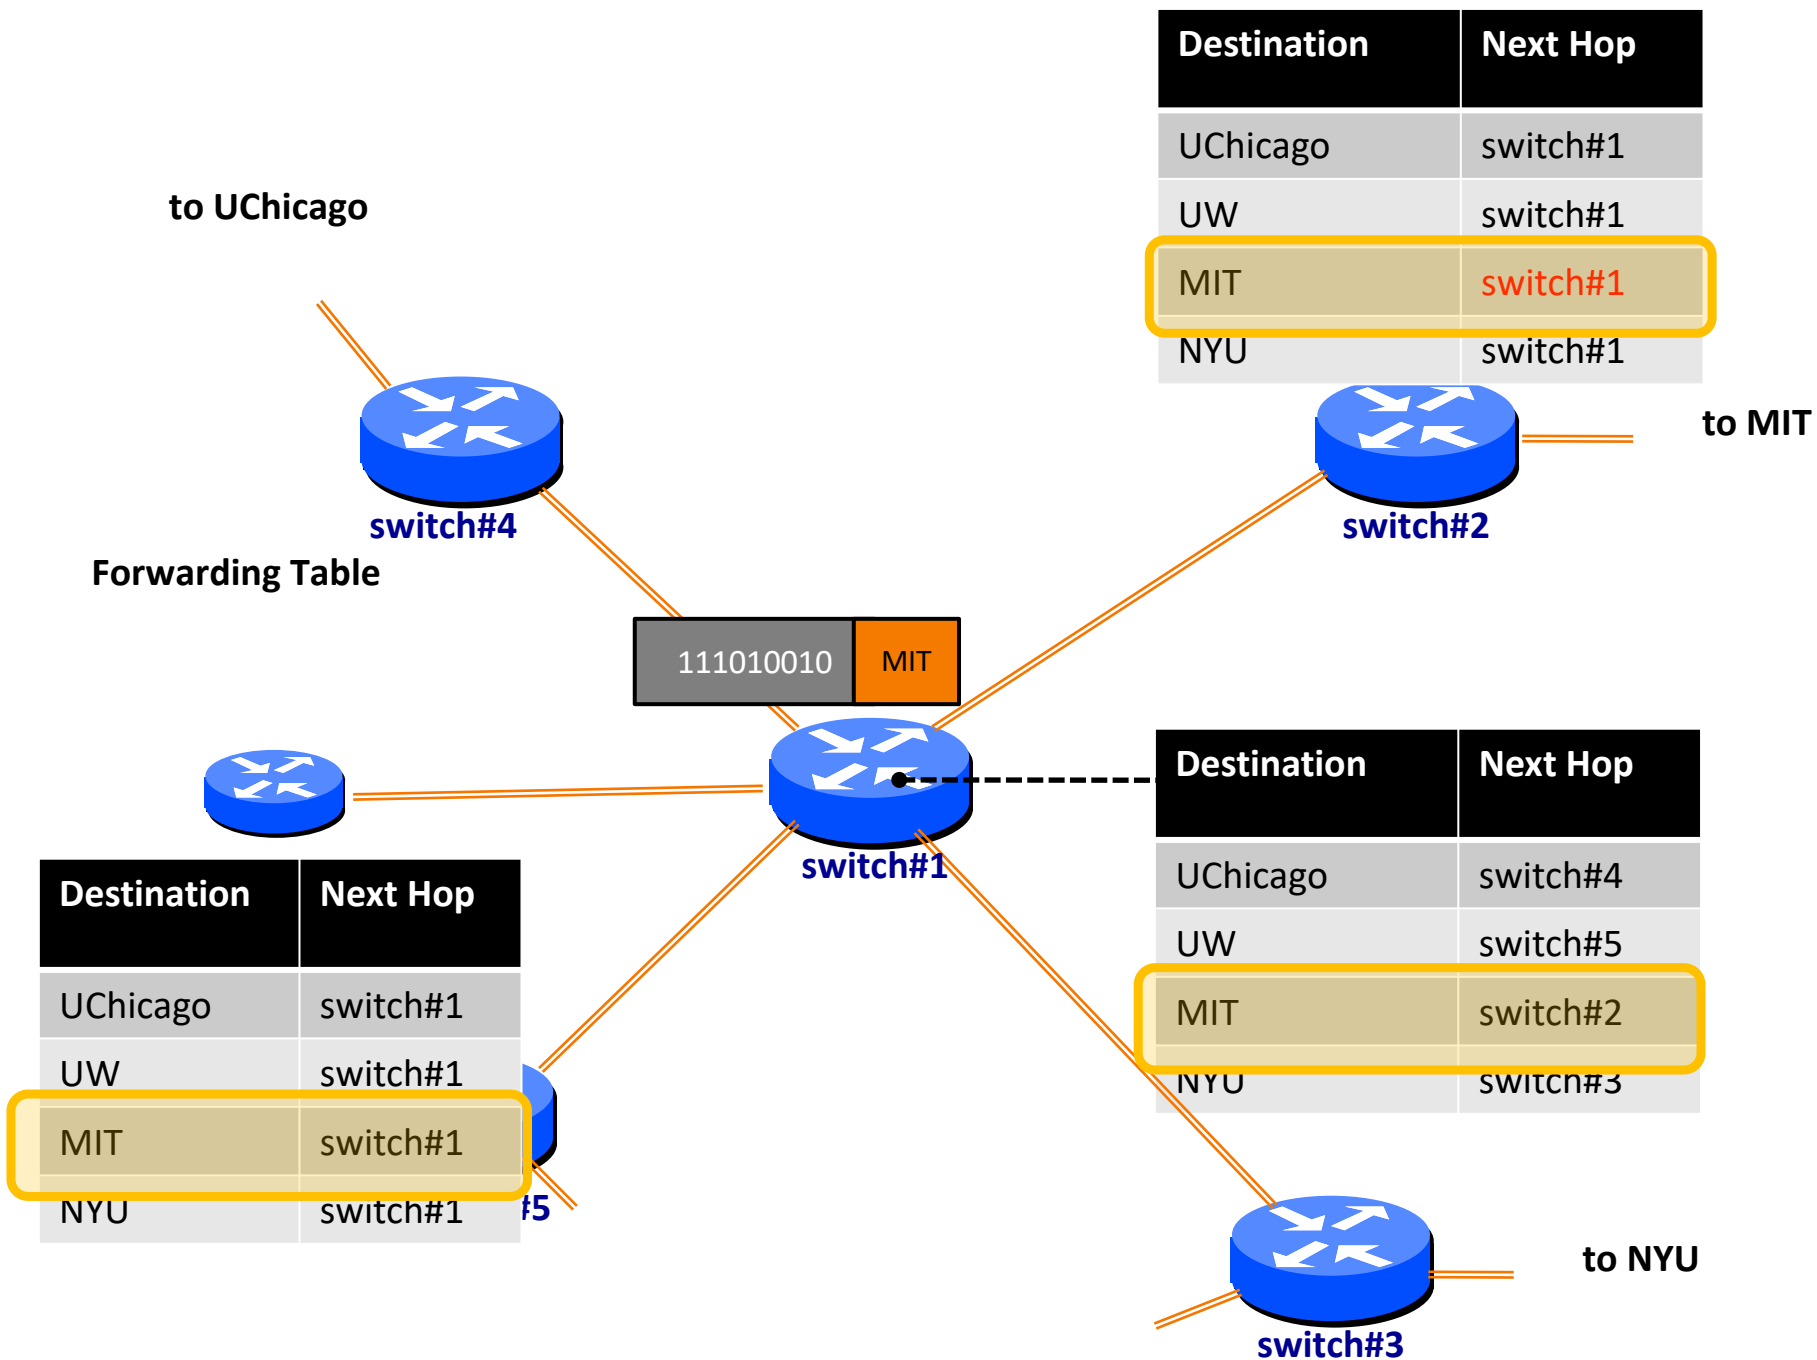
- **Validity** of routing state

Goal (v1)

- ▶ Find a path to a given destination
- ▶ How do we know that the state contained in forwarding tables meets our goal?
 - *this is what “validity” of routing state tells us*
 - *[this is non-standard terminology]*

Local vs. Global View of State

- *Local* routing state is the forwarding table in a single router
 - By itself, the state in a single router can't be evaluated
 - It must be evaluated in terms of the global context



Local vs. Global View of State

- *Local* routing state is the forwarding table in a single router
 - By itself, the state in a single router can't be evaluated
 - It must be evaluated in terms of the global context
- *Global* state refers to the collection of forwarding tables in each of the routers
 - Global state determines which paths packets take

(Will discuss later where this routing state comes from)

“Valid” Routing State

- Global state is “valid” if it produces forwarding decisions that always deliver packets to their destinations
- Goal of routing protocols: compute valid state
 - But how can you tell if routing state is valid?
- Need a succinct correctness condition for routing
 - Suggestions?

Necessary and Sufficient Condition

- Global routing state is valid *if and only if*:
 - There are no dead ends (other than destination)
 - There are no loops
- A dead end is when there is no outgoing link (next-hop)
 - A packet arrives, but the forwarding decision does not yield any outgoing link
- A loop is when a packet cycles around the same set of nodes forever



Loop!

to UChicago



switch#4

Forwarding Table



switch#1




switch#2

to MIT

Destination	Next Hop
UChicago	switch#1
UW	switch#1
MIT	switch#1
NYU	switch#1

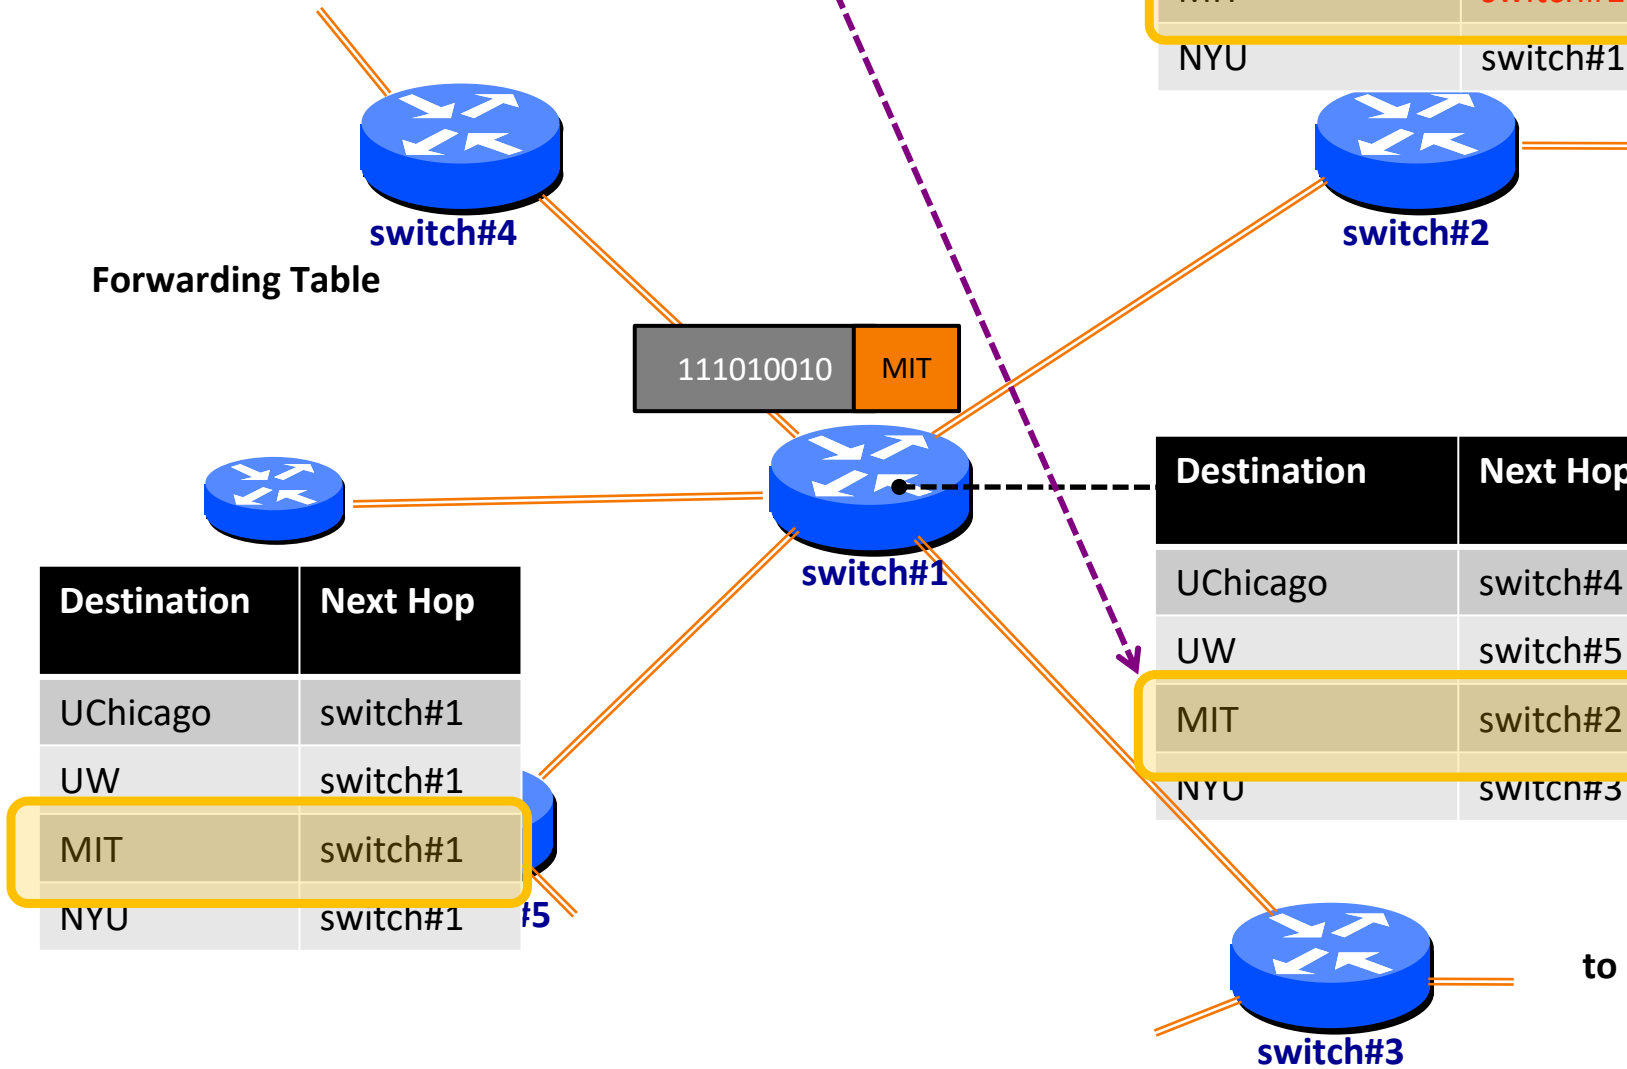
Destination	Next Hop
UChicago	switch#4
UW	switch#5
MIT	switch#2
NYU	switch#3

Destination	Next Hop
UChicago	switch#1
UW	switch#1
MIT	switch#1
NYU	switch#1



switch#3

to NYU



Deadend
(to MIT)

Destination	Next Hop
UChicago	switch#1
UW	switch#1
MIT	switch#1
NYU	switch#1

to UChicago

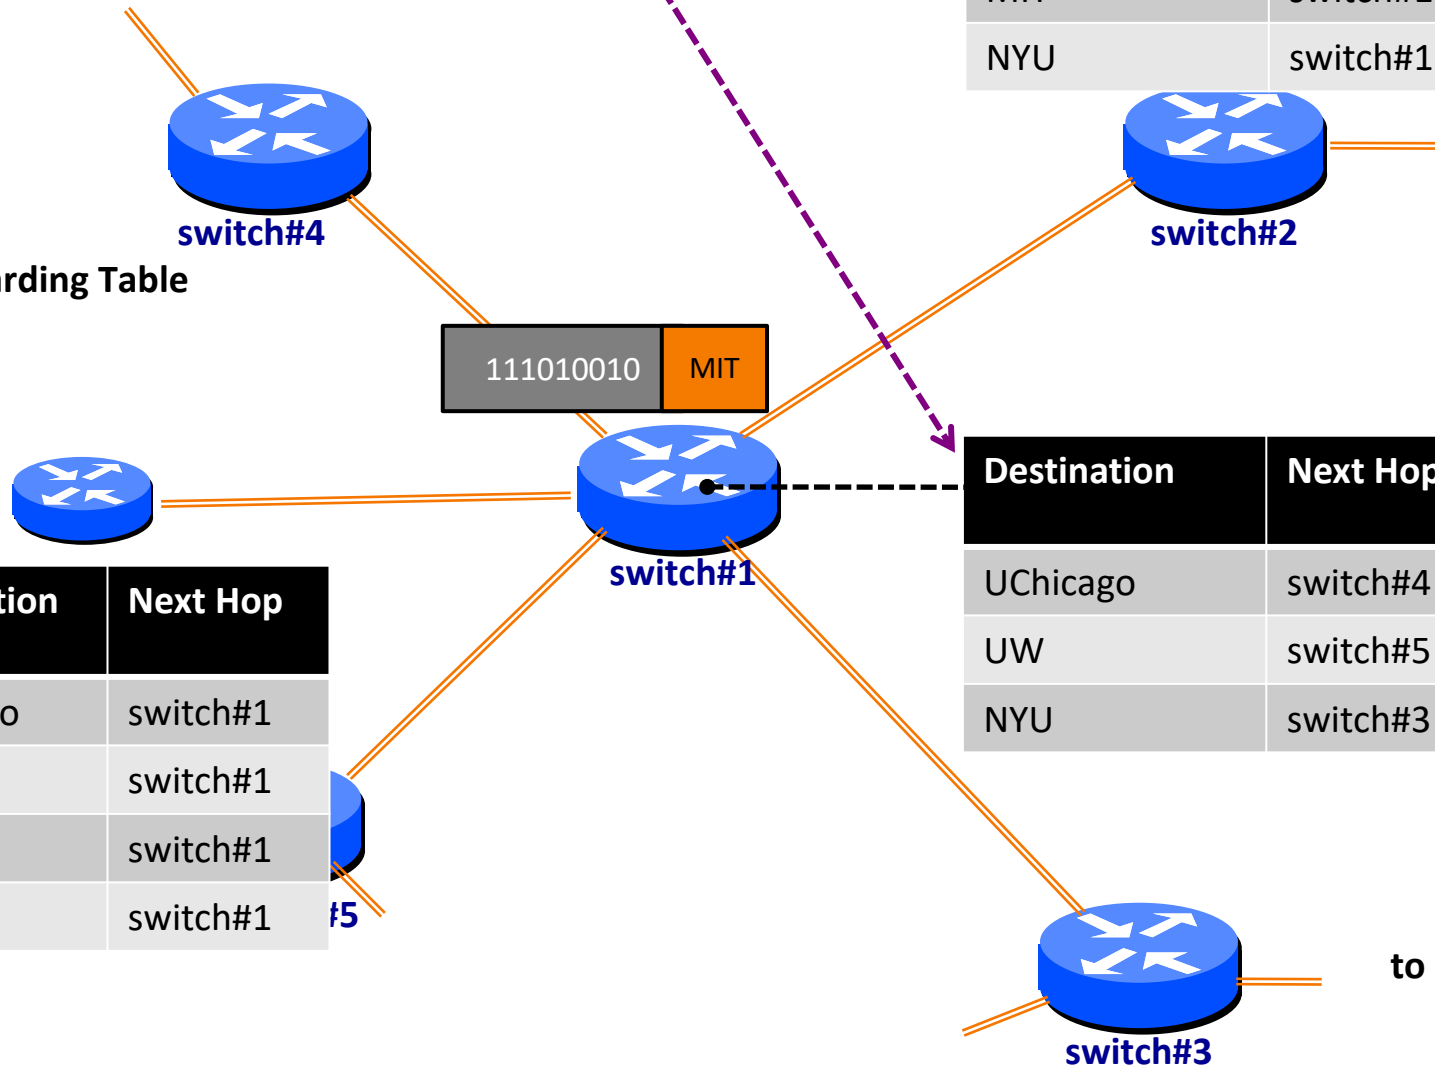
Forwarding Table

to MIT

Destination	Next Hop
UChicago	switch#1
UW	switch#1
MIT	switch#1
NYU	switch#1

Destination	Next Hop
UChicago	switch#4
UW	switch#5
NYU	switch#3

to NYU



Necessary and Sufficient Condition

- Global routing state is valid *if and only if*:
 - There are no dead ends (other than destination)
 - There are no loops

Necessary (“only if”): Easy

- If you run into a deadend before hitting destination, you’ll never reach the destination
- If you run into a loop, you’ll never reach destination


Sufficient (“if”):

- Assume no deadends, no loops
- Packet must keep wandering, but without repeating
 - If ever enter same switch from same link, will loop
- Only a finite number of possible links for it to visit
 - It cannot keep wandering forever without looping
 - Must eventually hit destination


Checking Validity of Routing State

- Focus only on a single destination
 - Ignore all other routing state
- Mark outgoing link (“next hop”) with arrow
 - There is only one at each node
- Eliminate all links with no arrows
- Look at what’s left....


Example 1




Pick Destination



Put arrows on outgoing links (to green dot)




Remove Unused Links




Leaves Spanning Tree: Valid

Second Example



Second Example



Is this valid?

Lesson....

- Very easy to check validity of routing state for a particular destination
- Deadends are obvious
 - Node without outgoing arrow
- Loops are obvious
 - Disconnected from rest of graph


Goal (for routing)

v1: Find a path to a given destination

v2: Find a *least cost path* to a given destination


Routing

- Goal: determine a “good” path through the network from source to destination
 - Good means usually the shortest path
- Network modeled as a graph
 - Routers \rightarrow nodes
 - Link \rightarrow edges
 - o Edge cost: delay, congestion level,...




Routing Problem

- Assume
 - A network with N nodes, where each edge is associated a cost
 - A node knows **only** its neighbors and the cost to reach them
- How does each node learn how to reach every other node along the shortest path?




Distance Vector: Control Traffic

- When the routing table of a node changes, it sends table to neighbors
 - A node updates its table with information received from neighbors



Example: Distance Vector Algorithm



Node A

Dest.	Cost	NextHop
B	2	B
C	7	C
D	∞	-

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	3	D

1 Initialization:

```

2  for all neighbors V do
3    if V adjacent to A
4       $D(A, V) = c(A, V);$ 
5    else
6       $D(A, V) = \infty;$ 

```

...


Node C

Dest.	Cost	NextHop
A	7	A
B	1	B
D	1	D

Node D

Dest.	Cost	NextHop
A	∞	-
B	3	B
C	1	C

Example: 1st Iteration ($C \rightarrow A$)



Node A

Dest.	Cost	NextHop
B	2	B
C	7	C
D	∞	-

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	3	D

$(D(C,A), D(C,B), D(C,D))$

Node C

Dest.	Cost	NextHop
A	7	A
B	1	B
D	1	D


Node D

Dest.	Cost	NextHop
A	∞	-
B	3	B
C	1	C

```

...
7 loop:
...
12 else if (update  $D(V, Y)$  received from  $V$ )
13   for all destinations  $Y$  do
14     if (destination  $Y$  through  $V$ )
15        $D(A, Y) = D(A, V) + D(V, Y)$ ;
16     else
17        $D(A, Y) = \min(D(A, Y),$ 
                         $D(A, V) + D(V, Y));$ 
18   if (there is a new minimum for dest.  $Y$ )
19     send  $D(A, Y)$  to all neighbors
20 forever
  
```

Example: 1st Iteration (C → A)



Node A

Dest.	Cost	NextHop
B	2	B
C	7	C
D	8	C

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	3	D

$$D(A,D) = \min(D(A,D), D(A,C) + D(C,D)) \\ = \min(\infty, 7 + 1) = 8$$

Node C

Dest.	Cost	NextHop
A	7	A
B	1	B
D	1	D


Node D

Dest.	Cost	NextHop
A	∞	-
B	3	B
C	1	C

```

...
7 loop:
...
12 else if (update D(V, Y) received from V)
13   for all destinations Y do
14     if (destination Y through V)
15       D(A, Y) = D(A, V) + D(V, Y);
16     else
17       D(A, Y) = min(D(A, Y),
18                     D(A, V) + D(V, Y));
19   if (there is a new minimum for dest. Y)
20     send D(A, Y) to all neighbors
21 forever
    
```

Example: 1st Iteration (B → A)



Node A

Dest.	Cost	NextHop
B	2	B
C	7	C
D	8	C

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	3	D



...

7 loop:

...

12 **else if** (update $D(V, Y)$ received from V)13 **for all** destinations Y **do**14 **if** (destination Y through V)15 $D(A, Y) = D(A, V) + D(V, Y);$ 16 **else**17 $D(A, Y) = \min(D(A, Y),$
 $D(A, V) + D(V, Y));$ 18 **if** (there is a new minimum for dest. Y)19 **send** $D(A, Y)$ to all neighbors20 **forever**


Node C

Dest.	Cost	NextHop
A	7	A
B	1	B
D	1	D

Node D

Dest.	Cost	NextHop
A	∞	-
B	3	B
C	1	C

Example: 1st Iteration ($B \rightarrow A$, $C \rightarrow A$)



Node A

Dest.	Cost	NextHop
B	2	B
C	3	B
D	5	B

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	3	D

$$D(A,D) = \min(D(A,D), D(A,B) + D(B,D)) \\ = \min(8, 2 + 3) = 5$$

$$D(A,C) = \min(D(A,C), D(A,B) + D(B,C)) \\ = \min(7, 2 + 1) = 3$$

...
7 loop:

```

...
12 else if (update D(V, Y) received from V)
13   for all destinations Y do
14     if (destination Y through V)
15       D(A, Y) = D(A, V) + D(V, Y);
16     else
17       D(A, Y) = min(D(A, Y),
18                     D(A, V) + D(V, Y));
19   if (there is a new minimum for dest. Y)
20     send D(A, Y) to all neighbors
21 forever
  
```


Node C

Dest.	Cost	NextHop
A	7	A
B	1	B
D	1	D

Node D

Dest.	Cost	NextHop
A	∞	-
B	3	B
C	1	C

Example: End of 1st Iteration



Node A

Dest.	Cost	NextHop
B	2	B
C	3	B
D	5	B

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	2	C

Node C

Dest.	Cost	NextHop
A	3	B
B	1	B
D	1	D


Node D

Dest.	Cost	NextHop
A	4	B
B	3	B
C	1	C

```

...
7 loop:
...
12 else if (update D(V, Y) received from V)
13   for all destinations Y do
14     if (destination Y through V)
15       D(A, Y) = D(A, V) + D(V, Y);
16     else
17       D(A, Y) = min(D(A, Y),
                     D(A, V) + D(V, Y));
18   if (there is a new minimum for dest. Y)
19     send D(A, Y) to all neighbors
20 forever
  
```

Example: End of 3rd Iteration



```

...
7 loop:
...
12 else if (update D(V, Y) received from V)
13   for all destinations Y do
14     if (destination Y through V)
15       D(A, Y) = D(A, V) + D(V, Y);
16     else
17       D(A, Y) = min(D(A, Y),
                     D(A, V) + D(V, Y));
18   if (there is a new minimum for dest. Y)
19     send D(A, Y) to all neighbors
20 forever
  
```

Node A

Dest.	Cost	NextHop
B	2	B
C	3	B
D	4	B

Node B

Dest.	Cost	NextHop
A	2	A
C	1	C
D	2	C

Node C

Dest.	Cost	NextHop
A	3	B
B	1	B
D	1	D

Node D

Dest.	Cost	NextHop
A	4	C
B	2	C
C	1	C

Nothing changes → algorithm terminates


0

Distance Vector: Link Cost Changes

```

7 loop:
8   wait (link cost update or update message)
9   if ( $c(A,V)$  changes by  $d$ )
10    for all destinations  $Y$  through  $V$  do
11       $D(A,Y) = D(A,Y) + d$ 
12    else if (update  $D(V, Y)$  received from  $V$ )
13      for all destinations  $Y$  do
14        if (destination  $Y$  through  $V$ )
15           $D(A,Y) = D(A,V) + D(V, Y)$ ;
16        else
17           $D(A, Y) = \min(D(A, Y), D(A, V) + D(V, Y))$ ;
18      if (there is a new minimum for destination  $Y$ )
19        send  $D(A, Y)$  to all neighbors
20  forever

```



Node	D	C	N
Node B	A	4	A
Node B	C	1	C
Node C	A	5	B
Node C	B	1	B

Node	D	C	N
Node B	A	1	A
Node B	C	1	C
Node C	A	5	B
Node C	B	1	B

Node	D	C	N
Node B	A	1	A
Node B	C	1	C
Node C	A	2	B
Node C	B	1	B

Node	D	C	N
Node B	A	1	A
Node B	C	1	C
Node C	A	2	B
Node C	B	1	B

“good news travels fast”

Link cost changes here

Algorithm terminates



time

DV: Count to Infinity Problem

```

7 loop:
8   wait (link cost update or update message)
9   if ( $c(A,V)$  changes by  $d$ )
10    for all destinations  $Y$  through  $V$  do
11       $D(A,Y) = D(A,Y) + d$ 
12    else if (update  $D(V, Y)$  received from  $V$ )
13      for all destinations  $Y$  do
14        if (destination  $Y$  through  $V$ )
15           $D(A,Y) = D(A,V) + D(V, Y);$ 
16        else
17           $D(A, Y) = \min(D(A, Y), D(A, V) + D(V, Y));$ 
18      if (there is a new minimum for destination  $Y$ )
19        send  $D(A, Y)$  to all neighbors
20  forever

```




Link cost changes here; note that B also maintains the shortest route to A for C, which is 6. Thus $D(B, A)$ becomes 6 !


time

Distance Vector: Poisoned Reverse

- If C routes through B to get to A:
 - C tells B its (C's) distance to A is infinite (so B won't route to A via C)
 - Will this completely solve count to infinity problem?




NO! multihop loops



Link State: Control Traffic

- Each node floods its local information to every other node in the network
- Each node ends up knowing the **entire** network topology → use Dijkstra to compute the shortest path to every other node



Link State: Node State

