

Decentralized / Peer to Peer Lecture 15

Today: Peer-to-Peer

- Unstructured Peer-to-Peer Protocols
 - How and why do they work?
 - Napster, Gnutella, Kazaa
 - BitTorrent, Freenet, Skype
- Structured Peer-to-Peer Overlays
 - A quick overview

Traditional Internet Services Model

- Client-server
 - Many clients, 1 (or more) server(s)
 - Web servers, DNS, file downloads, video streaming
- Problems
 - Scalability: how many users can a server support?
What happens when user traffic overload servers?
Limited resources (bandwidth, CPU, storage)
 - Reliability: if # of servers is small, what happens when they break, fail, get disconnected, are mismanaged by humans?
 - Efficiency: if your users are spread across the entire globe, how do you make sure you answer their requests quickly?

The Alternative: Peer-to-Peer

- A simple idea
 - Users bring their own resources to the table
 - A cooperative model: clients = peers = servers
- The benefits
 - Scalability: # of “servers” grows with users
BYOR: bring your own resources (storage, CPU, B/W)
 - Reliability: load spread across many peers, probability of them all failing is VERY low...
 - Efficiency: peers are distributed. Peers locate and get service from nearby peers (Is this always possible?)

The Peer-to-Peer Challenge

- What are the key components for leveraging P2P?
 - Communication: how do peers talk to each other
 - Service / data location: how do peers know who to talk to?
- Answers depend on the application context
 - The obvious example: file-sharing
 - Communication?
 - Service / data location?
- We'll discuss these in detail w.r.t. different protocols
 - Napster/Maze, Gnutella, Kazaa, BitTorrent, Freenet, Skype
 - Many others (Mojonation, Overnet, ICQ)

Unstructured P2P Applications

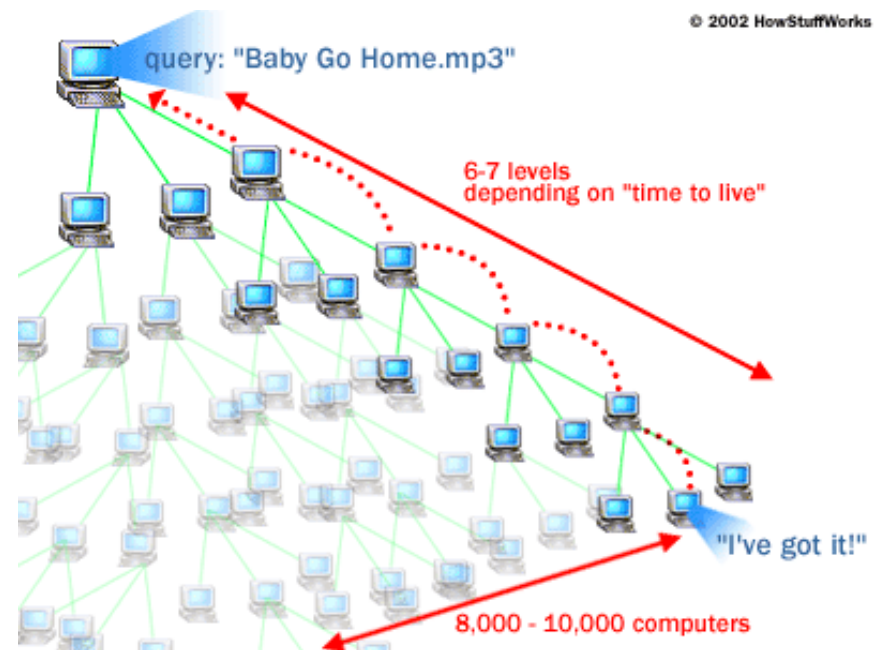
- Solves part of the problem (file location mostly)
 - Focus on locating popular objects
 - Hierarchy of super-nodes index client file collections
 - Queries sent as controlled flood amongst supernodes
- Limitations
 - File search
 - Does not support generic any-to-any node communication
 - Probabilistic
 - Can only search a small portion of files in system
otherwise would flood network with query traffic (e. g. Gnutella)
 - Uncommon files (in a file system) are easily lost
the more unique your file is, more likely it will be inaccessible

Centralized Approach


- The original: Napster
- Centralized index server(s)
 - Supported all queries
- What caused its downfall?
 - Not scalability
 - Centralization of liability

Gnutella

- The original Gnutella
 - Recent incarnations look very much like Kazaa
- Flat network
 - Join via bootstrap node
 - Connect to random set of existing hosts
 - Resolve queries by localized flooding
 - Time to live fields limit hops
- Problem
 - High bandwidth costs in control msgs
 - Flood of queries took up all available bandwidth




File Search via Flooding in Gnutella




Resilience to Failures and Attacks

- Previous studies (Barabasi) show interesting dichotomy of resilience for “scale-free networks”
 - Resilient to random failures, but not attacks
- Here’s what it looks like for Gnutella




1771 Peers in Feb, 2001




After random 30% of peers removed

Resilience to Failures and Attacks

- Previous studies (Barabasi) show interesting dichotomy of resilience for “scale-free networks”
 - Resilient to random failures, but not attacks
- Here’s what it looks like for Gnutella




1771 Peers in Feb, 2001



After top 4% of peers are removed

Hierarchical P2P Networks

- FastTrack network (Kazaa, Grokster, Morpheus, Gnutella++)



Kazaa

- Very popular from its inception
 - Hierarchical flooding helps improve scale
 - Large shift to broadband helped quite a bit as well
 - Based in Europe, more relaxed copyright laws
- Susceptible to poison attacks
 - Mainly used by RIAA-like organizations
 - Get online and distribute large # of popular files
 - Insert static or silence in songs to make them useless
 - Quite effective in dissuading downloaders

Onwards to Bigger Things...

- Kazaa works well for audio
 - User churn is still a problem
 - High enough to disrupt downloads of large files
- BitTorrent
 - Ideal for larger files, movies, dvds, large collections...
 - Leverages simultaneous downloads by many users
 - A “seed” has a complete copy, forwards different blocks to different users
 - Users then share blocks amongst themselves
 - Typically 256KB blocks, verified using SHA-1 checksums
 - Key is incentive model for downloads: Tit-for-Tat
 - Send blocks to those who you are downloading from
 - Optimistic unchoke to bring in new users

BitTorrent Operation

- Centralized boot-strap mechanism
 - Need a .torrent file containing length, name, hashes, URL of tracker
 - Easily locatable via HTTP servers (e.g. piratebay)
- The Tracker
 - Coordinator that matches up leechers with seeds & each other
 - Maintains info about all currently active clients
- What's new?
 - Trackerless BitTorrent clients (Azureus etc.)
 - Now use structured overlays (DHTs) instead

Outline

- Unstructured Peer-to-Peer Protocols
 - How and why do they work?
 - Napster, Gnutella, Kazaa
 - BitTorrent, Freenet, Skype
- Structured Peer-to-Peer Overlays
 - Quick overview of protocols
 - Applications

Why Do We Need Structure?

- Without structure, the search problem is harder
 - Anything can be anywhere
 - Need to look through many (or all) nodes to have confidence in finding object X
- So how do we add structure? Ans: Add rules
 - Give everyone (every machine) a name
 - Give every object a name
 - Match up all object names to machine names
 - If you are looking for X, $\text{mapping}(X) \rightarrow Y$, talk to node Y
- Challenges
 - This mapping must be easy, embedded into the network
 - It must be consistent as the machines in network change
 - One new machine should not disrupt everything

Structured Overlays


- Application infrastructure for large-scale apps
- Two key pieces of functionality
 - Scalable any to any communication in network of size N
 - Much larger scale compared to distance vector/link state
 - Each machine only needs to know $\text{Log}(N)$ other nodes
 - Routing from $A \rightarrow B$ takes at most $\text{Log}(N)$ steps or hops
 - Deterministic key-node mapping
 - Allows peer rendezvous using common name
 - This is the object to node mapping we need
 - We call it Key-Based Routing

Peer-to-Peer Protocols

- Unstructured Peer to Peer Approaches
 - Napster, Gnutella, KaZaa
 - ❖ probabilistic search (optimized for the hay, not the needle)
 - ❖ locality-agnostic routing (resulting in high network b/w costs)
- Structured Peer to Peer Overlays
 - the first protocols (2001): Tapestry, Pastry, Chord, CAN
 - then: Kademlia, SkipNet, Viceroy, Symphony, Koorde, Ulysseus...
 - ❖ **distinction**: how to choose your neighbors
 - ❖ Tapestry/Chimera, Pastry: latency-optimized routing mesh
 - ❖ **distinction**: application interface
 - ❖ distributed hash table: put (key, data); data = get (key);
 - ❖ Tapestry/Chimera: decentralized object location and routing

Structured Overlays at 10000 ft

- Node IDs and keys from randomized namespace (SHA-1)
 - incremental routing towards destination ID
 - each node has small set of outgoing routes, e.g. prefix routing
 - $\log(n)$ neighbors per node, $\log(n)$ hops between any node pair




The Common Denominator

- Key-based Routing layer (Tier 0)
 - Large sparse Id space N (160 bits: $0 - 2^{160}$ represented as base b)
 - Nodes in overlay network have $\text{nodeIds} \in N$
 - Given $k \in N$, overlay deterministically maps k to its root node (a live node in the network)
- Intuition
 - Take set of all possible names
 - Divide space “evenly” into “buckets” stored at live nodes in network
 - Provide pointers between nodes so nodes can find desired bucket
- Main call: `route (key, msg)`

Chord

- NodeIDs are numbers on ring
- Closeness defined by numerical proximity
- Finger table
 - keep routes for next node 2^i away in namespace
 - routing table size: $\log_2 n$
 - n = total # of nodes
- Routing
 - iterative hops from source
 - at most $\log_2 n$ hops

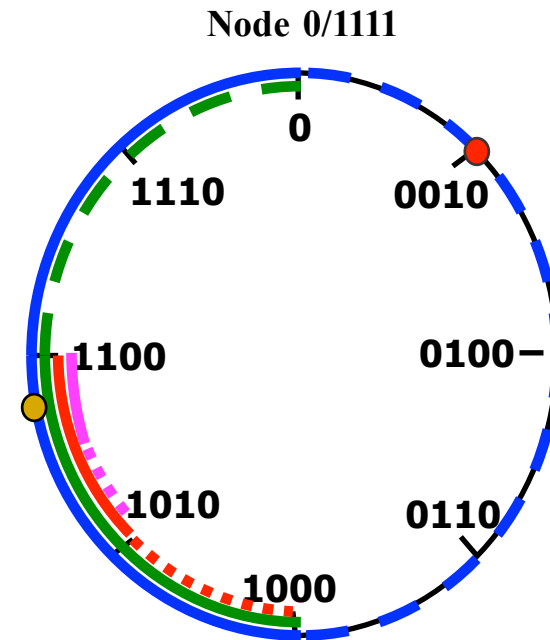


Questions About Chord

- Why is it $\log_2 N$ hops?
- How does a node join the network?
- What is the “region” each node is responsible for?
- What happens when nodes fail?



Tapestry / Pastry / Chimera

- incremental prefix routing
 - $0010 \rightarrow 1XXX \rightarrow 10XX \rightarrow 101X \rightarrow 1011$
- routing table
 - keep nodes matching at least i digits to destination
 - table size: $b * \log_b n$
- routing
 - recursive routing from source
 - at most $\log_b n$ hops



Routing in Detail

Example: Octal digits, 2^{12} namespace, $2175 \rightarrow 0157$



Today: Peer-to-Peer

- Unstructured Peer-to-Peer Protocols
 - How and why do they work?
 - Napster, Gnutella, Kazaa
 - BitTorrent, Freenet, Skype
- Structured Peer-to-Peer Overlays
 - A quick overview