

Final Review

Lecture 16

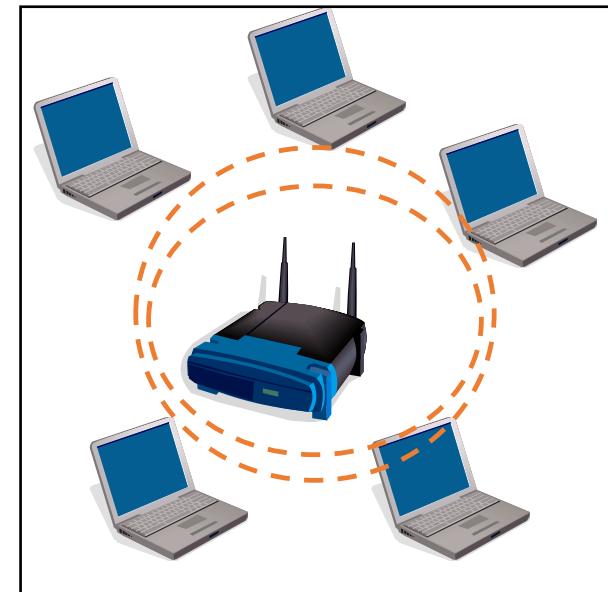
Material List (after Mid-term 1)

- Lecture 11: Wireless networking
- Lecture 12: Wireless Routing & TCP
- Lecture 13: Datacenter Networking
- Lecture 14: Data-driven Design (TCP Remi)
- Lecture 15: P2P Networks

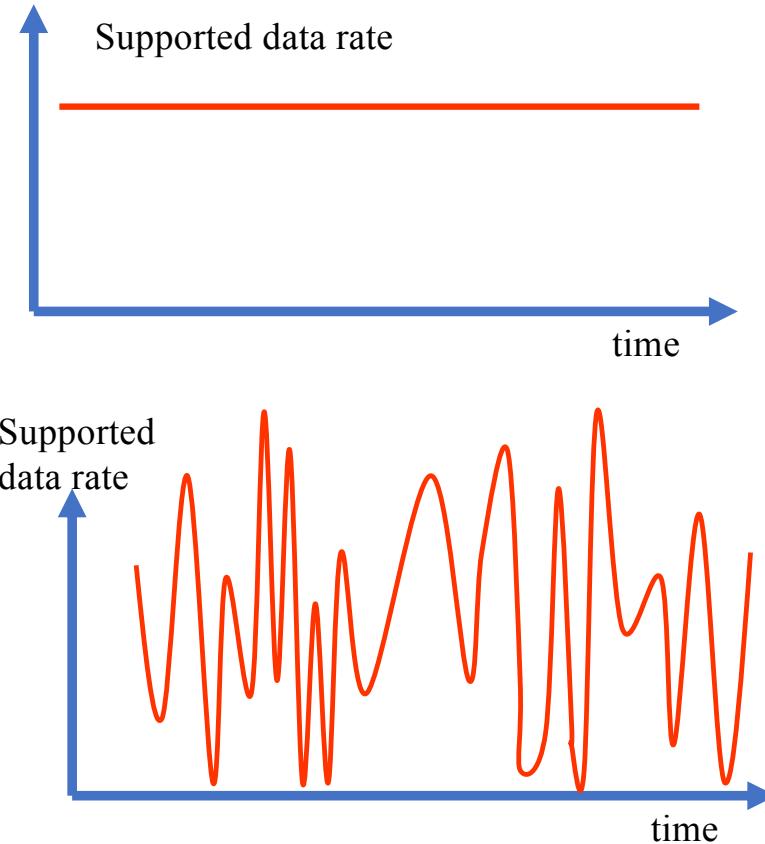
Wireless

Wireless Networking

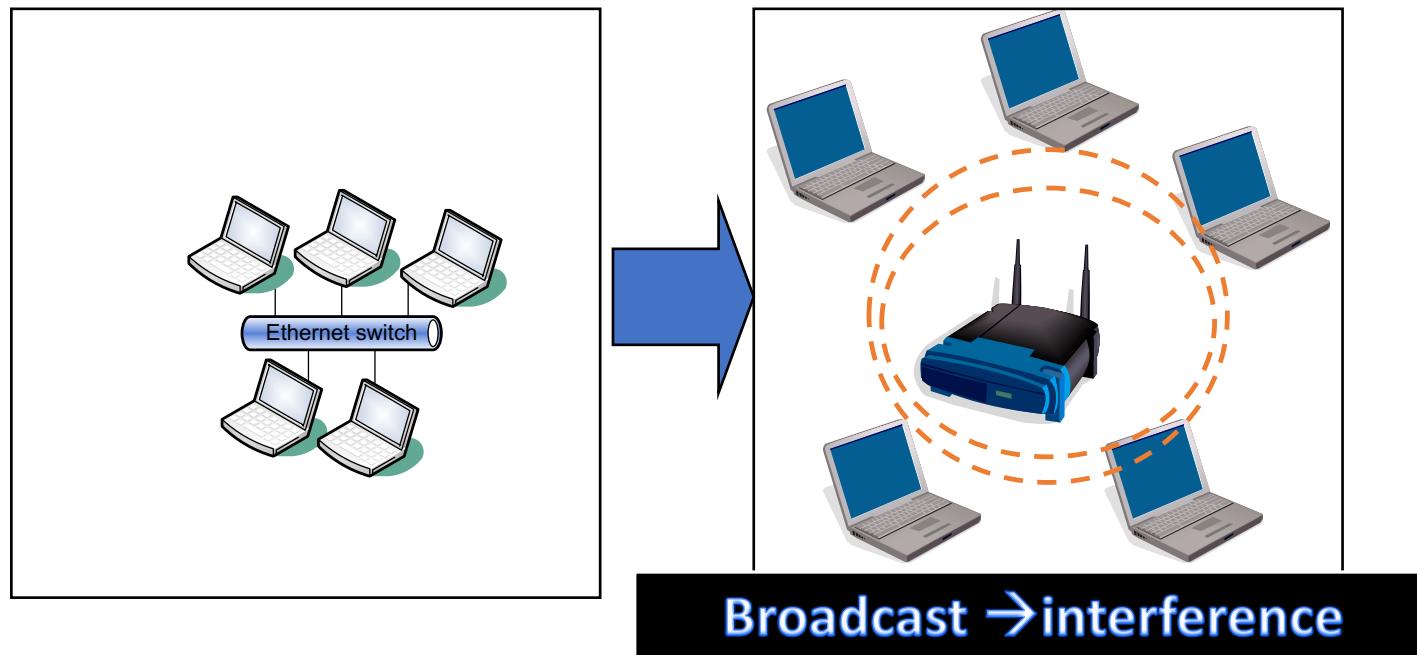
- Changes in Physical Medium
- Key artifacts
 - Link fluctuations
 - Interference



The Difference: # 1



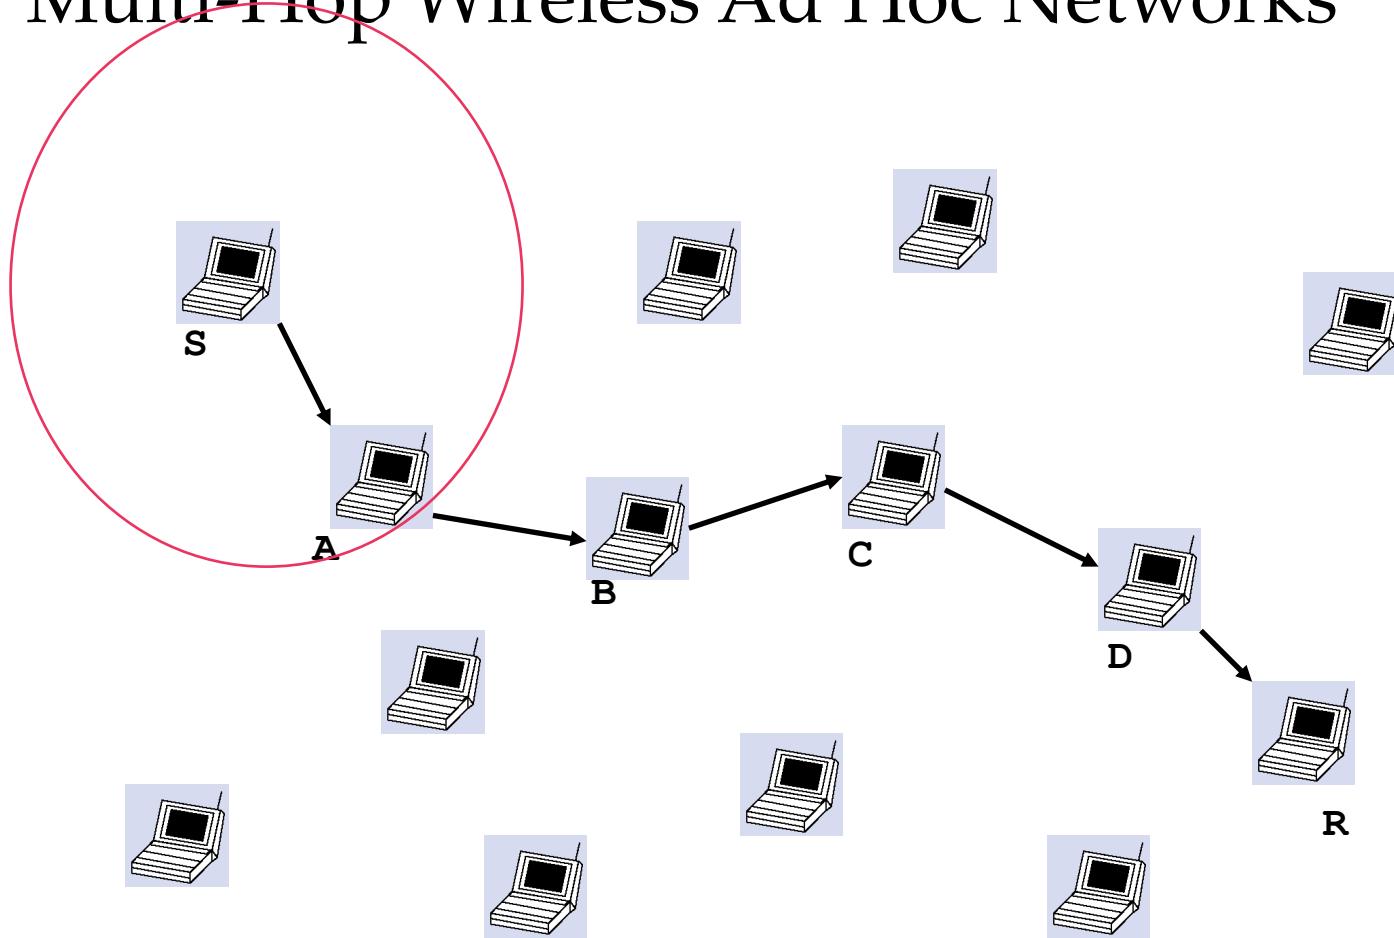
The Difference #2



Wireless Networking Solutions

- LINK Layer: 802.11 CSMA/CA
 - The problems of Hidden terminal & Exposed terminal
 - What triggered them (roles of transmitter & receiver)
 - How to address them
 - What is CSMA, what is CA
 - How to avoid interference using channel assignment
- Networking Layer: Multi-hop routing
 - Why wireless multi-hop routing is harder/less efficient than wired multi-hop routing
 - What really happens (RTS/CTS/transmission)
 - TCP makes it even harder, why?

Multi-Hop Wireless Ad Hoc Networks



Courtesy of Tianbo Kuang and Carey Williamson University of Calgary)

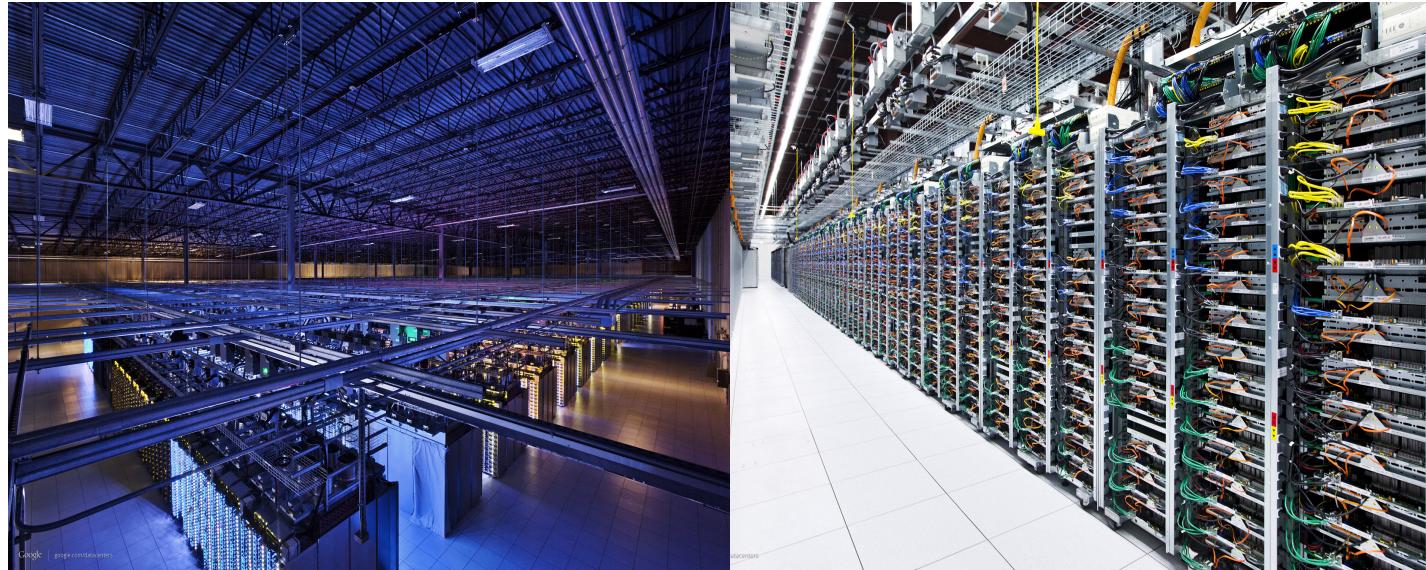
Wireless Networking Solutions

- Transport Layer: Wireless TCP
 - Running TCP over wireless.. What is the problem, how does it differ from running TCP over just wired networks
 - 2 groups of strategies to address wireless TCP
 - Masking
 - Notification
 - Where are they implemented? Sender, Receiver, intermediate nodes?
 - Good and bad of these approaches

Data Center

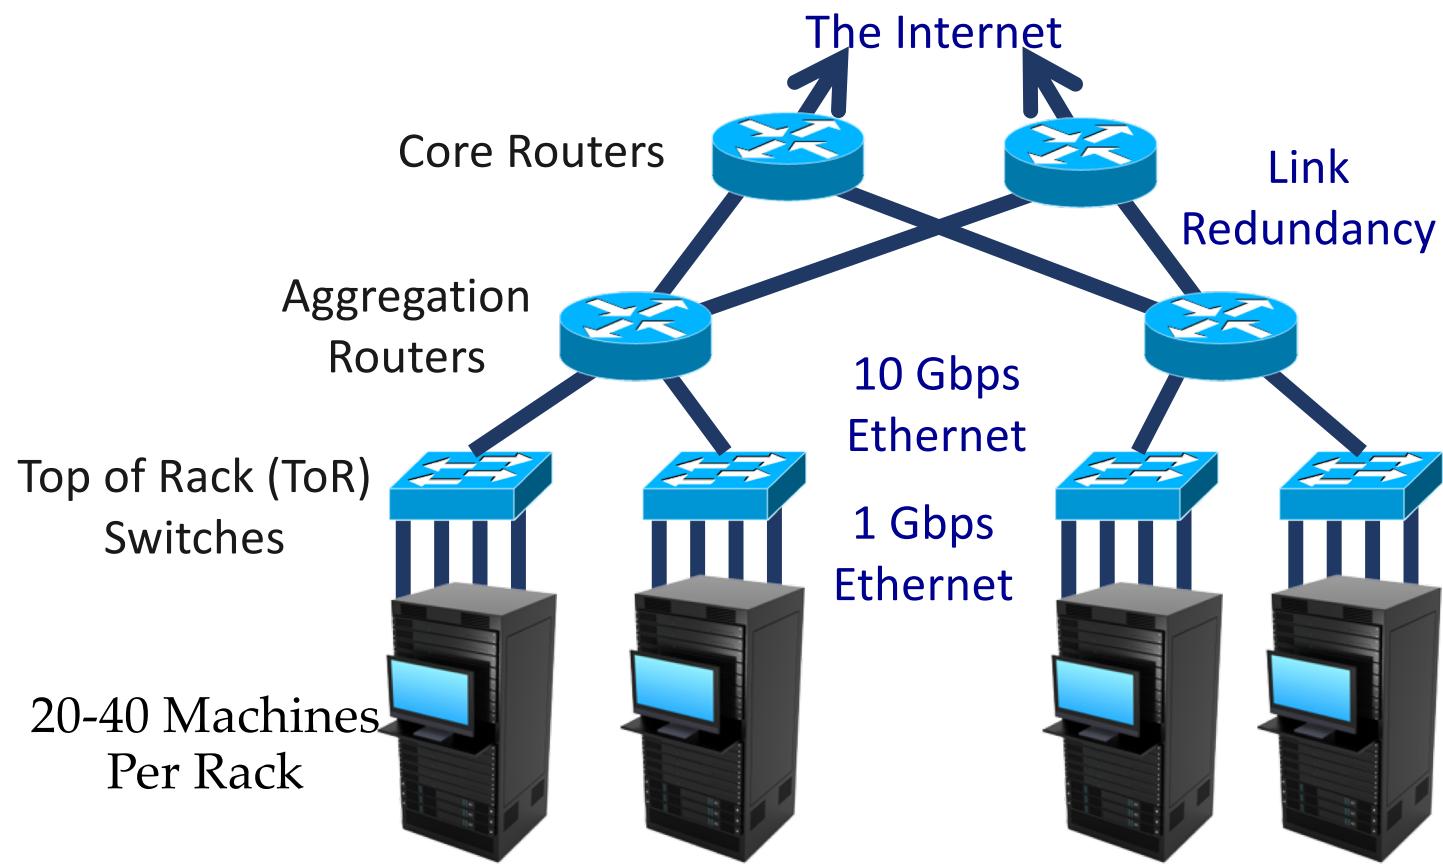
Today's Data Centers

- Warehouse of servers

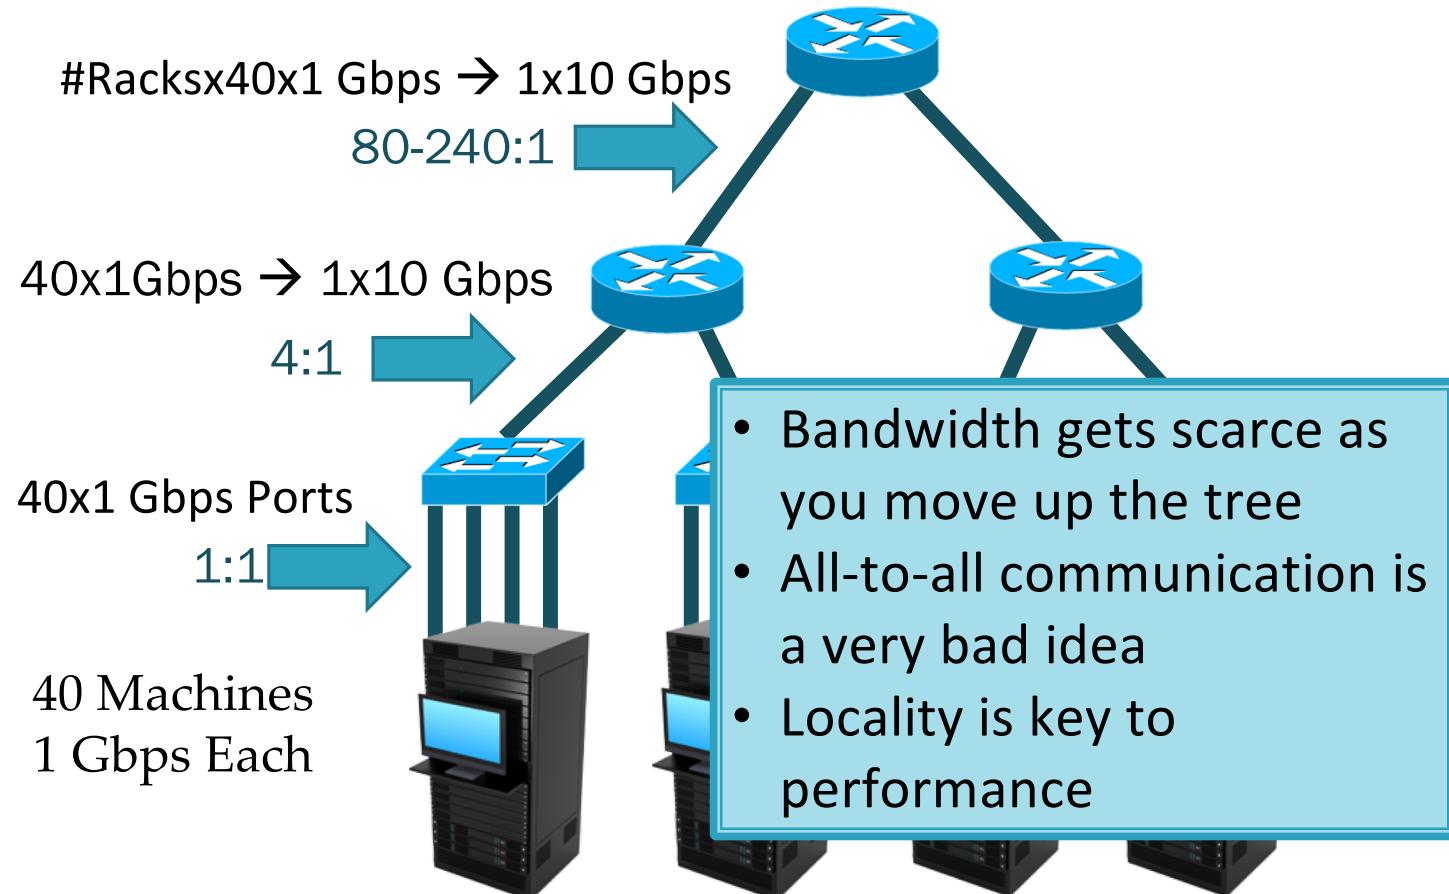


<https://www.youtube.com/watch?v=zDAYZU4A3w0>

Typical Data Center Topology



Problem: Oversubscription



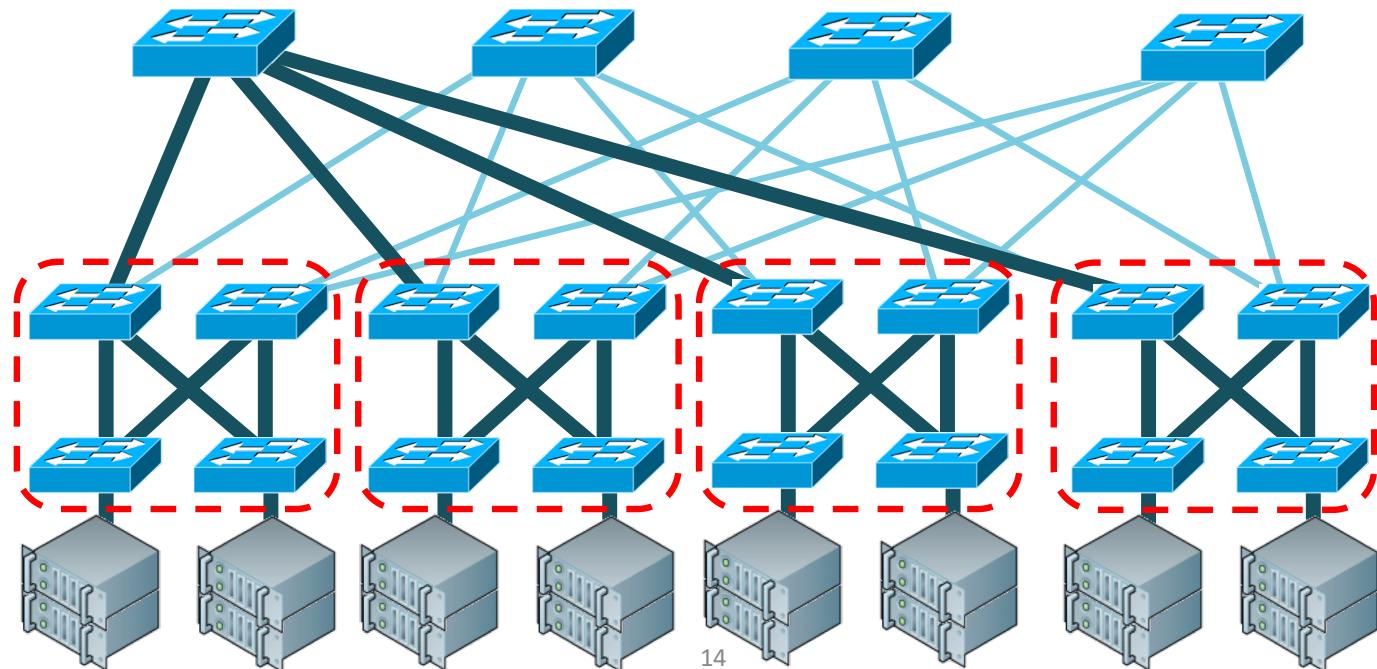
Fat-tree

To build a K-ary fat tree

- K-port switches
- K pods, each with K switches
- $K^3/4$ hosts
- $(K/2)^2$ core switches

In this example K=4

- 4-port switches
- 4 pods, each with 4 switches
- $K^3/4 = 16$ hosts
- $(K/2)^2 = 4$ core switches



Summary

- Network configuration in data center
 - Why it differs from Internet
- What is the problem of oversubscription
- How do we solve the problem of oversubscription
 - Fat tree
 - How do you build fat tree
 - What happens if one router breaks (Hint: Facebook video)
- What is the problem of traffic hotspot in Data center
 - What are the solutions (wired solution vs. wireless solution) to address this problem

Summary (cont)

- TCP in Data center
 - Difference in requirements
 - Partition/aggregation pattern
 - Delay sensitive
 - Elephant flow vs. mice flow
 - Can change sender/receiver since they are controlled by the admin
 - Dirty slate approach: DCTCP
 - Adapting existing TCP
 - Explicit notification
 - Adjust window in proportional to congestion
 - Clean slate approach: D³
 - Completely modify TCP to perform delay-based resource reservation
 - How to implement this??

Data driven Design

Data-Driven Network Design

- Type 1: Understand “real” network behaviors from “in the wild” measurements
 - Many existing network protocol designs were driven by insights (which might be obsolete)
- Type 2: Use machine learning, especially reinforcement learning to identify rules for network functionalities

TCP Remy

- So far, human designers have created TCP's congestion-control algorithms **by hand**
 - Making implicit assumptions
- The **emergent behavior** of a flock of machines running TCP is complicated and difficult to characterize
- Remy: **machine generated** TCP algorithm based on observations of the current network state

<http://web.mit.edu/remy/>

Remy: a program that generates
congestion-control schemes offline

Input:

- ▶ Prior assumptions (what network may do)
- ▶ Goal (what app wants)

Output: CC algorithm for a TCP sender (RemyCC)

Time: a few hours

Cost: \$5–\$10 on Amazon EC²

Remy Structure

- Remy searches for the best congestion-control algorithm
- Optimizes expected objective over prior assumptions
- Makes tractable by limiting available state
- Remy tracks three congestion signals

r_ewma: moving average of interval between acks

s_ewma: ... between sender timestamps echoed in acks

rtt_ratio: ratio of last RTT to smallest RTT so far

Remy maps each state to an action

$$\text{RULE}(r_{\text{ewma}}, s_{\text{ewma}}, \text{rtt_ratio}) \rightarrow \langle m, b, \tau \rangle$$

m Multiple to congestion window

b Increment to congestion window

τ Minimum interval between two outgoing packets

On ack:

- ▶ $\langle m, b, \tau \rangle \leftarrow \text{RULE}(r_{\text{ewma}}, s_{\text{ewma}}, \text{rtt_ratio})$
- ▶ $\text{cwnd} \leftarrow m \cdot \text{cwnd} + b$

Send packet if:

- ▶ $\text{cwnd} > \text{FlightSize}$, and
- ▶ last packet sent $> \tau$ ago

Things that need extra care

- Remy assumes that the designer can model the network explicitly, i.e. making some assumptions on network model
 - What about mismatched set of assumptions

P2P

Summary

- How P2P differs from traditional client/server architecture
- Unstructured P2P
 - Design for ???
 - What does “unstructured” mean?
 - Which one is easy to find in unstructured P2P: popular items vs. non-popular items, why?
- Structured P2P
 - What is the structure for? Why do we need structure
 -

Why Do We Need Structure?

- Without structure, the search problem is harder
 - Anything can be anywhere
 - Need to look through many (or all) nodes to have confidence in finding object X
- So how do we add structure? Ans: Add rules
 - Give everyone (every machine) a name
 - Give every object a name
 - Match up all object names to machine names
 - If you are looking for X, mapping(X)→Y, talk to node Y
- Challenges
 - This mapping must be easy, embedded into the network
 - It must be consistent as the machines in network change
 - One new machine should not disrupt everything

Structured Overlays

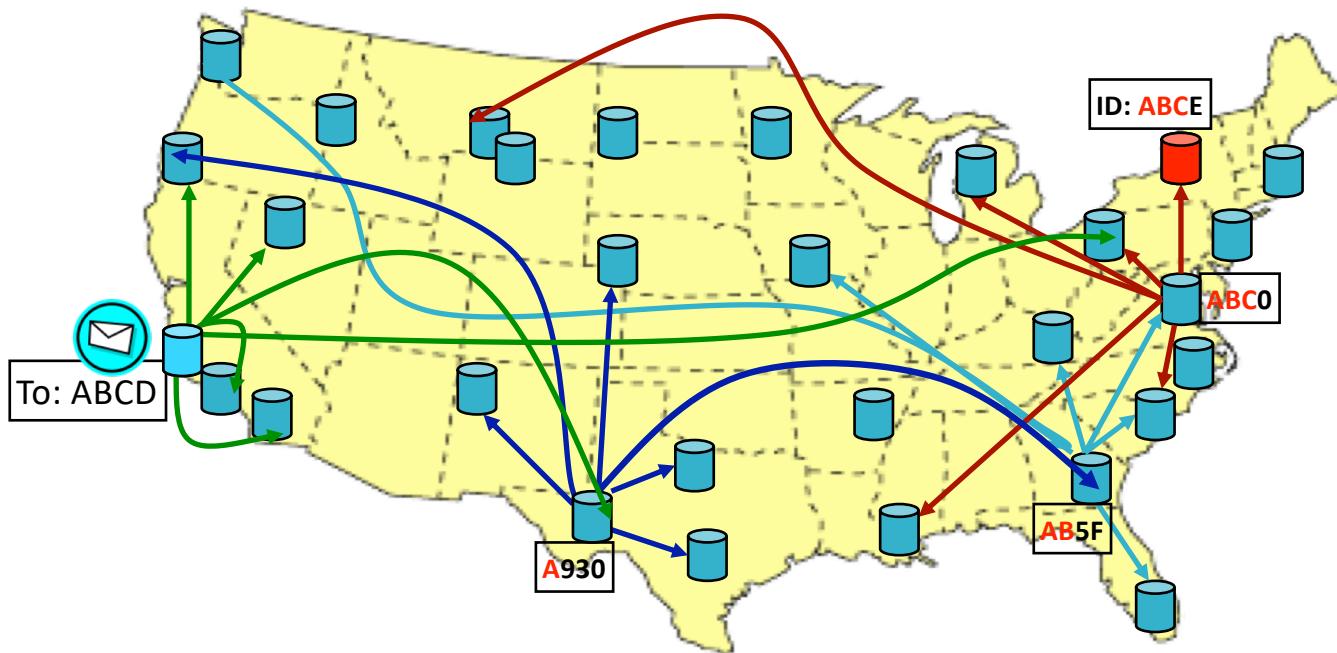
- Application infrastructure for large-scale apps
- Two key pieces of functionality
 - Scalable any to any communication in network of size N
 - Much larger scale compared to distance vector/link state
 - Each machine only needs to know $\text{Log}(N)$ other nodes
 - Routing from A->B takes at most $\text{Log}(N)$ steps or hops
 - Deterministic key-node mapping
 - Allows peer rendezvous using common name
 - This is the object to node mapping we need
 - We call it Key-Based Routing

Peer-to-Peer Protocols

- Unstructured Peer to Peer Approaches
 - Napster, Gnutella, KaZaa
 - ❖probabilistic search (optimized for the hay, not the needle)
 - ❖locality-agnostic routing (resulting in high network b/w costs)
- Structured Peer to Peer Overlays
 - the first protocols (2001): Tapestry, Pastry, Chord, CAN
 - then: Kademlia, SkipNet, Viceroy, Symphony, Koorde, Ulysseus...
 - ❖**distinction:** how to choose your neighbors
 - ❖Tapestry/Chimera, Pastry: latency-optimized routing mesh
 - ❖**distinction:** application interface
 - ❖distributed hash table: put (key, data); data = get (key);
 - ❖Tapestry/Chimera: decentralized object location and routing

Structured Overlays at 10000 ft

- Node IDs and keys from randomized namespace (SHA-1)
 - incremental routing towards destination ID
 - each node has small set of outgoing routes, e.g. prefix routing
 - $\log(n)$ neighbors per node, $\log(n)$ hops between any node pair



This is it!

- Lecture 11: Wireless networking
- Lecture 12: Wireless Routing & TCP
- Lecture 13: Datacenter Networking
- Lecture 14: Data-driven Design (TCP Remi)
- Lecture 15: P2P Networks

Q&A Session