



Lecture 2 (Networking concepts)



Switched networks enable efficient scaling
Sharing:

Reservations -> circuit switch

src sends reservations request to dst
Switches establish circuit
src starts sending data
src sends teardown message



Pros: predictable performance, fast switching
Cons: circuit setup/teardown, bad with bursty traffic, if switch fails circuit will fail

On demand -> packet switch

Each packet has dst and treated independently, buffers absorb transient overload

Pros: efficient use of network resources, simple to implement, robust against switch fails

Cons: unpredictable performance, buffer management, congestion control

Overview: Communication via Sockets

- Client
 1. Create a socket with the `socket()` system call
 2. Bind the socket to an address using the `bind()` system call
 3. Send and receive data. The simplest is to use the `read()` and `write()` system calls.

- Server
 1. Create a socket with the `socket()` system call
 2. Bind the socket to an address using the `bind()` system call.
 3. Listen for connections with the `listen()` system call
 4. Accept a connection with the `accept()` system call (blocks until a client connects with the server)
 5. Send and receive data

UDP does not need listen, accept, connect

Lecture 3 (Performance metrics)

Delay: How long does it take to send a packet from dst to src?

• Consists of four components

- transmission delay
 - propagation delay
 - queuing delay
 - processing delay
- due to link properties
due to traffic mix and switch internals

Link bandwidth: bits/sec

Propagation delay: secs for 1 bit to move through link

BDP=bandwidth x propagation delay

1mb = 1,000,000; 1ms = 1/10³ s

Transmission delay: how long it takes to push all the bits of a packet into a link, packet size/link bandwidth

Queueing delay: how long a packet has to sit in buffer before being processed.

If arrival rate > departure rate, QD=inf.

A = average arrival rate

W = average packet wait time

L = average length of queue


L = A x W (hard to compute)

Processing delay: how long the switch takes to process a packet, negligible

Loss: fraction of packets dropped

Throughput: rate at which dst receiving packets, data size/transfer time


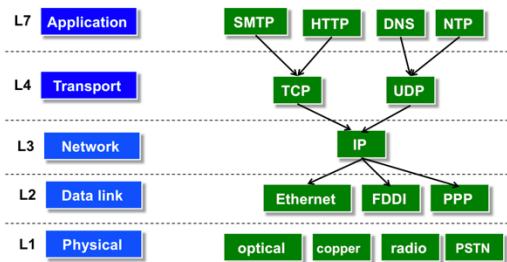
transmission rate R transmission rate R' > R



Transfer time = F/R + propagation delay + L/R'

Average throughput = min { R, R' } = R


Lecture 4 (IP routing algorithms)



Host – implement all 5 layers

Switches – L1 and L2

Routers – L1, L2, L3



Forwarding: Local router determines output link, read address from packet's network layer header, search forwarding table

Routing: Network wide process to determine content of forwarding tables (end to end path for each destinations)

Local routing is the forwarding table in a single router, global state refers to collection of forwarding tables in each of the routers

Routing state is valid if and only if there are no dead ends and there are no loops

Distance vector: only know neighbors, when routing table changes, send to neighbors

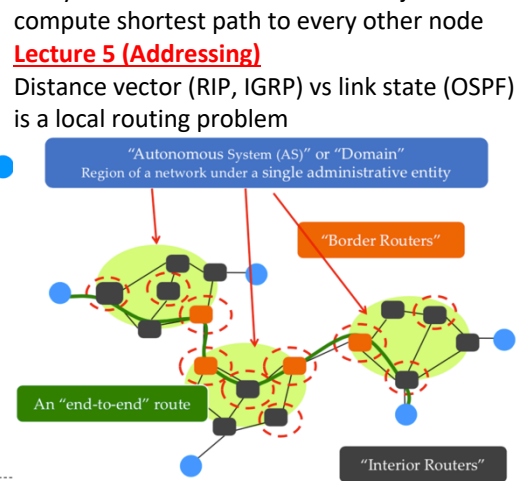
Distance vector algorithm, broadcast and receive neighboring forwarding tables, can calculate next hop for a destination and total cost to get to that destination. Good news travels fast and bad news travels slowly.

Poisoned reverse: distance is inf. If routes through itself

Link state: node floods local information to every other node in the network. Dijkstra to compute shortest path to every other node

Lecture 5 (Addressing)

Distance vector (RIP, IGRP) vs link state (OSPF) is a local routing problem



Forwarding table should have an entry for a range of addresses. Host addressing is key

AS wants policy, autonomy, privacy

BGP (Border Gateway Protocol) extends distance vector to accommodate policy

IPv4 is 32 bit number split into prefix (network) and suffix (host)

Classful addressing only has 3 sizes, 1 byte, 2 byte, 3 byte network prefixes

CIDR (classless Interdomain routing) – flexible division between network and host addresses.

128.23.9/26 means 26 bits as network prefix.

26 represented by 32 bit mask

ICANN -> ARIN -> ISP -> individuals

Lecture 6 (BGP)

AS can be customer, provider, peer

Peers have similar traffic and route through each other to save money

BGP selects best route based on policy rather than shortest distance; sends entire path for each destination (rather than just distance metric, avoid loops); selective route


advertisement; aggregate routes for different prefixes

eBGP – BGP between border routers

iBGP – between routers within same AS


IGP – intradomain routing protocol

Security: AS can claim to serve a prefix they do not have access to, AS can forward along a route different than advertised



With Gao-Rexford, the AS policy graph is a DAG (directed acyclic graph) and routes are "valley free"

Gao-Rexford guaranteed to converge



Persistent oscillation, AS path length mislead

Lecture 7, 8(IP, TCP)

IP Datagram – 20 bytes; Address resolution protocol ARP maps IP address to link level address to be used in direct delivery

Internet control message protocol (ICMP) – exchange control/error messages about delivery of IP datagrams: address mask, timestamp, source quench (traffic overload), parameter problem (errors in IP datagram field), echo, time exceeded (report expired datagrams), redirect, destination unreachable, traceroute (send datagrams each with increasing TTL)

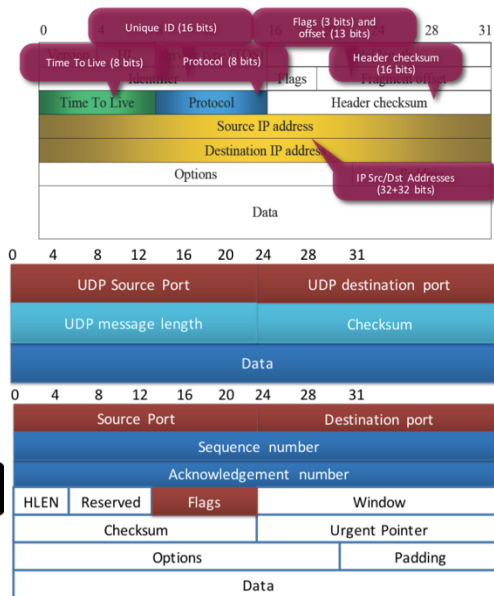
TCP: sequence # specifies position of the segment data, ack # specifies position of next byte expected from communication partner. **TCP flow control window:** make sure receiving end can handle data (no regard for network). **Accept iff $ack \# < syn \# < ack \# + window$** SYN/FIN – establishing/terminating connect ACK – when ACK is valid; URG – urgent pointer says where non urgent starts; PUSH; ABORT – abort connection.

NACK – packet did not arrive; cum. ACK – ACKS for all $k < n$; SACK – selective ACK

Key func: flow control, data transfer, congestion control, connection setup

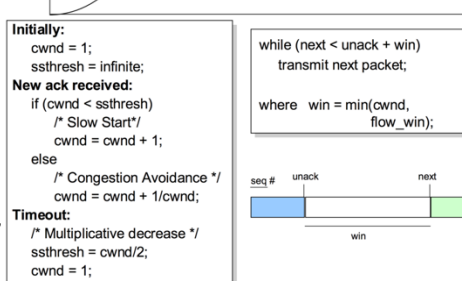

Throughput=window/RTT

Know packet dropped when time out, receive duplicate ACKs




TCP Congestion Control: guessing about state of network, additive increase multiplicative decrease, KNEE – point at which throughput increases slow and delay increases fast, CLIFF – point after which throughput decreases to zero fast. Congestion control stay left of cliff, congestion avoidance stay left of knee.

AIMD – converges to efficiency and fairness Maintain cwnd, flow_win, Ssthresh (update cwnd), for sending use $\min(cwnd, flow_win)$



MIAD – not stable not fair; AIAD – stable not fair; MIMD – stable not fair; AIMD – stable fair



Lecture 9 (FTP, DNS)



FTP: transfer files between 2 computers; need to resolve OS, character, naming, directory 2 connections: Control (persistent connection, 21), Data (ephemeral connection, 20)

In active FTP the client tells the server which port the server should connect to (server initiates connection). Client firewall would usually block this so we use passive FTP

DNS – local servers and resolver software

Active Mode (PORT)

Passive Mode (PASV)



Lecture 10 (DNS, RPC)

DNS caching: DNS servers cache responses to queries, delete cached entry after TTL expires, top level servers rarely change

Attacking DNS: impersonate local DNS server and give wrong IP address to client; denial of service the root to make them unavailable DNS properties: easy unique naming, fate sharing for network failures, reasonable trust model, caching lends scalability, performance

Server designs: iterative vs concurrent; connections vs connectionless; stateful vs stateless

Remote procedure call: called procedure not in same address space, transport independence, hides network from programmer, BLOCKS, there is more delay, cannot pass pointers, **idempotent can be repeated safely non idempotent cannot** RPC calls can be lost/duplicated, at least once/zero or more (reply/no reply)

Request – transaction id, procedure id, credentials, verifier, params(XDR); REPLY – trans id, status, verifier, status, resultsXDR XDR – specify how data object encoded, client server need to agree on format

Concurrent, Connection-oriented Server

1. Create a TCP socket (s1)
2. Bind it to a port (INADDR_ANY)
3. Place the socket in passive mode (listen)
4. Accept connection (uses s1, returns new socket s2)
5. Fork slave process
 1. Close old socket (C-s1)
 2. Interact with the client using new socket (C-s2)
 3. Close new connection (C-s2)
 4. Exit
6. Close new socket (s2)
7. Go to (4) or...
8. Exit